

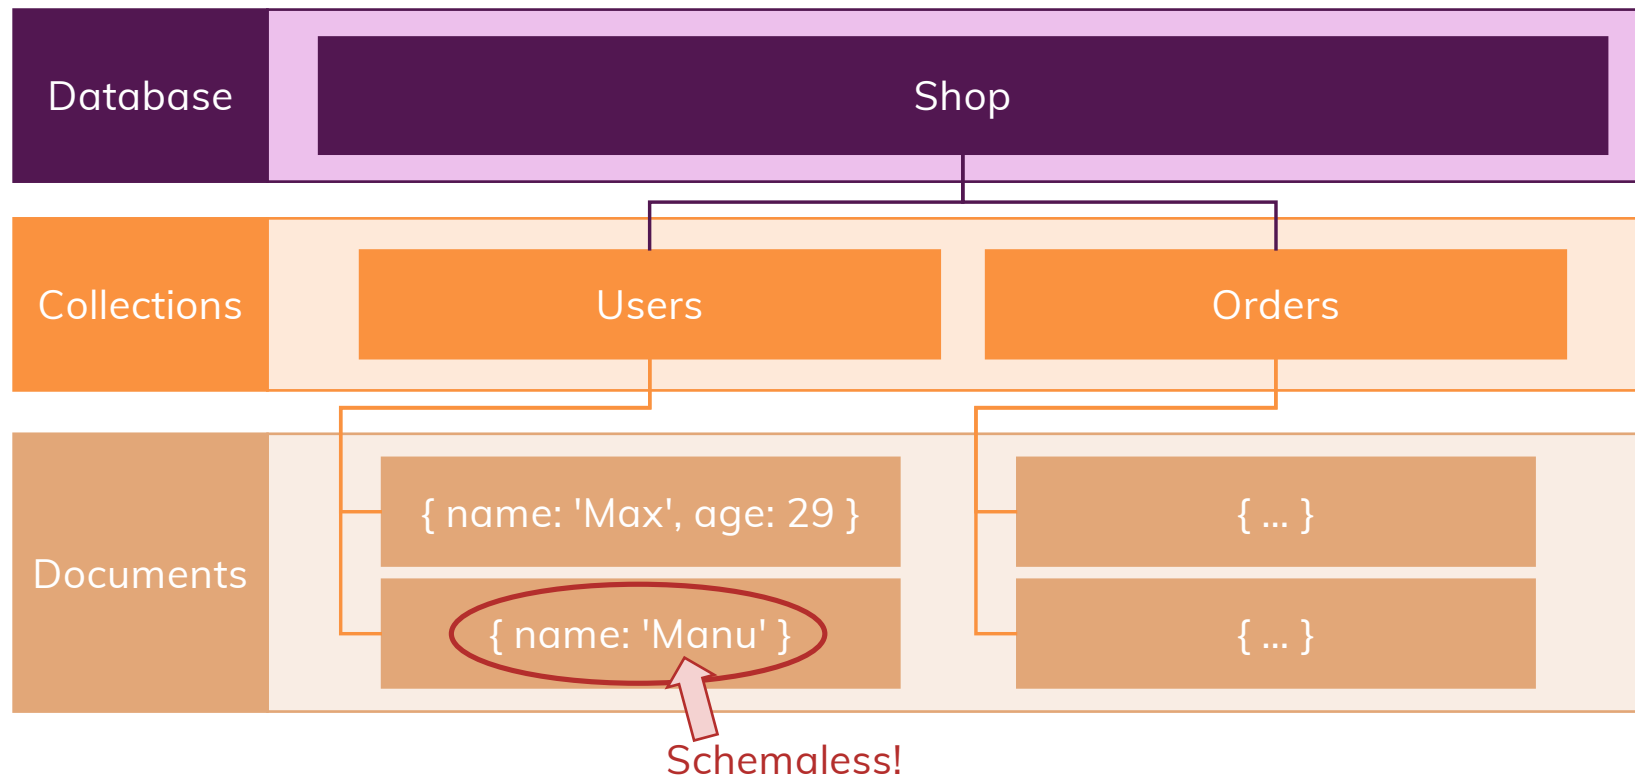
What?



Humongous

Because it can store lots and lots of data

How it works



JSON (BSON) Data Format

```
{  
  "name": "Max",  
  "age": 29,  
  "address":  
    {  
      "city": "Munich"  
    },  
  "hobbies": [  
    { "name": "Cooking" },  
    { "name": "Sports" }  
  ]  
}
```

BSON Data Structure

No Schema!

id: 1	"name": "Max"	"age": 29	...
id: 2	"name": "Manu"		...
id: 3		"age": 31	...

Users Collection

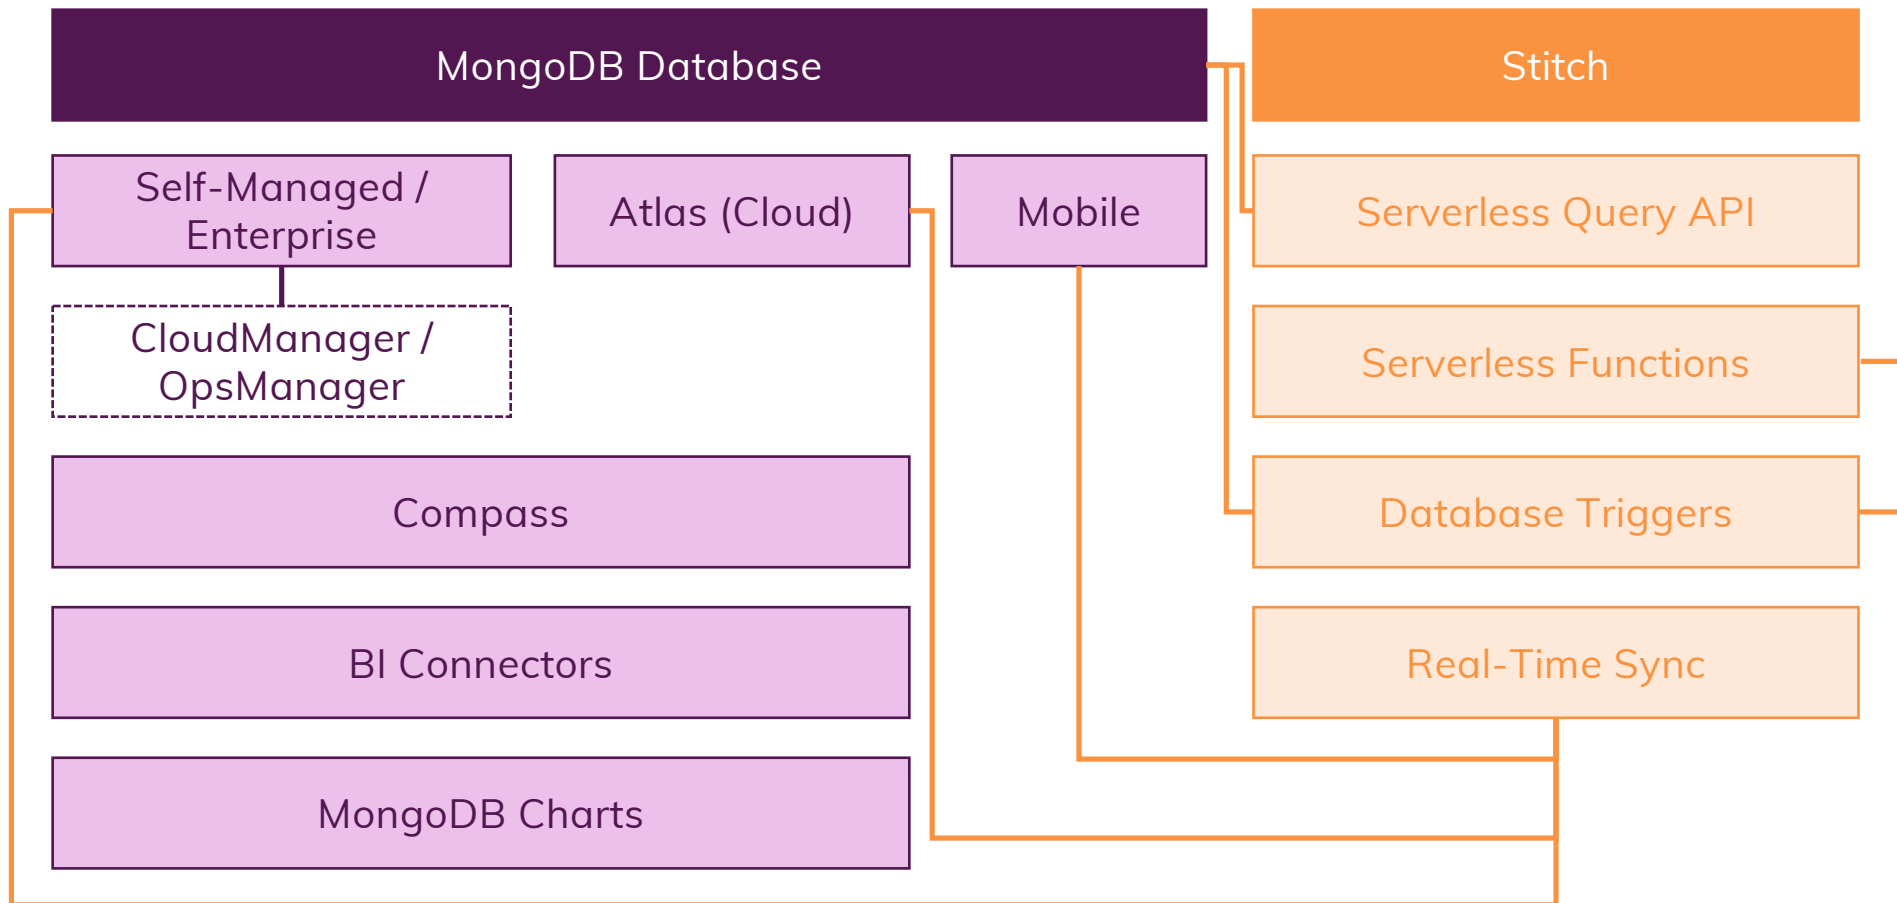
Relations

No / Few Relations!

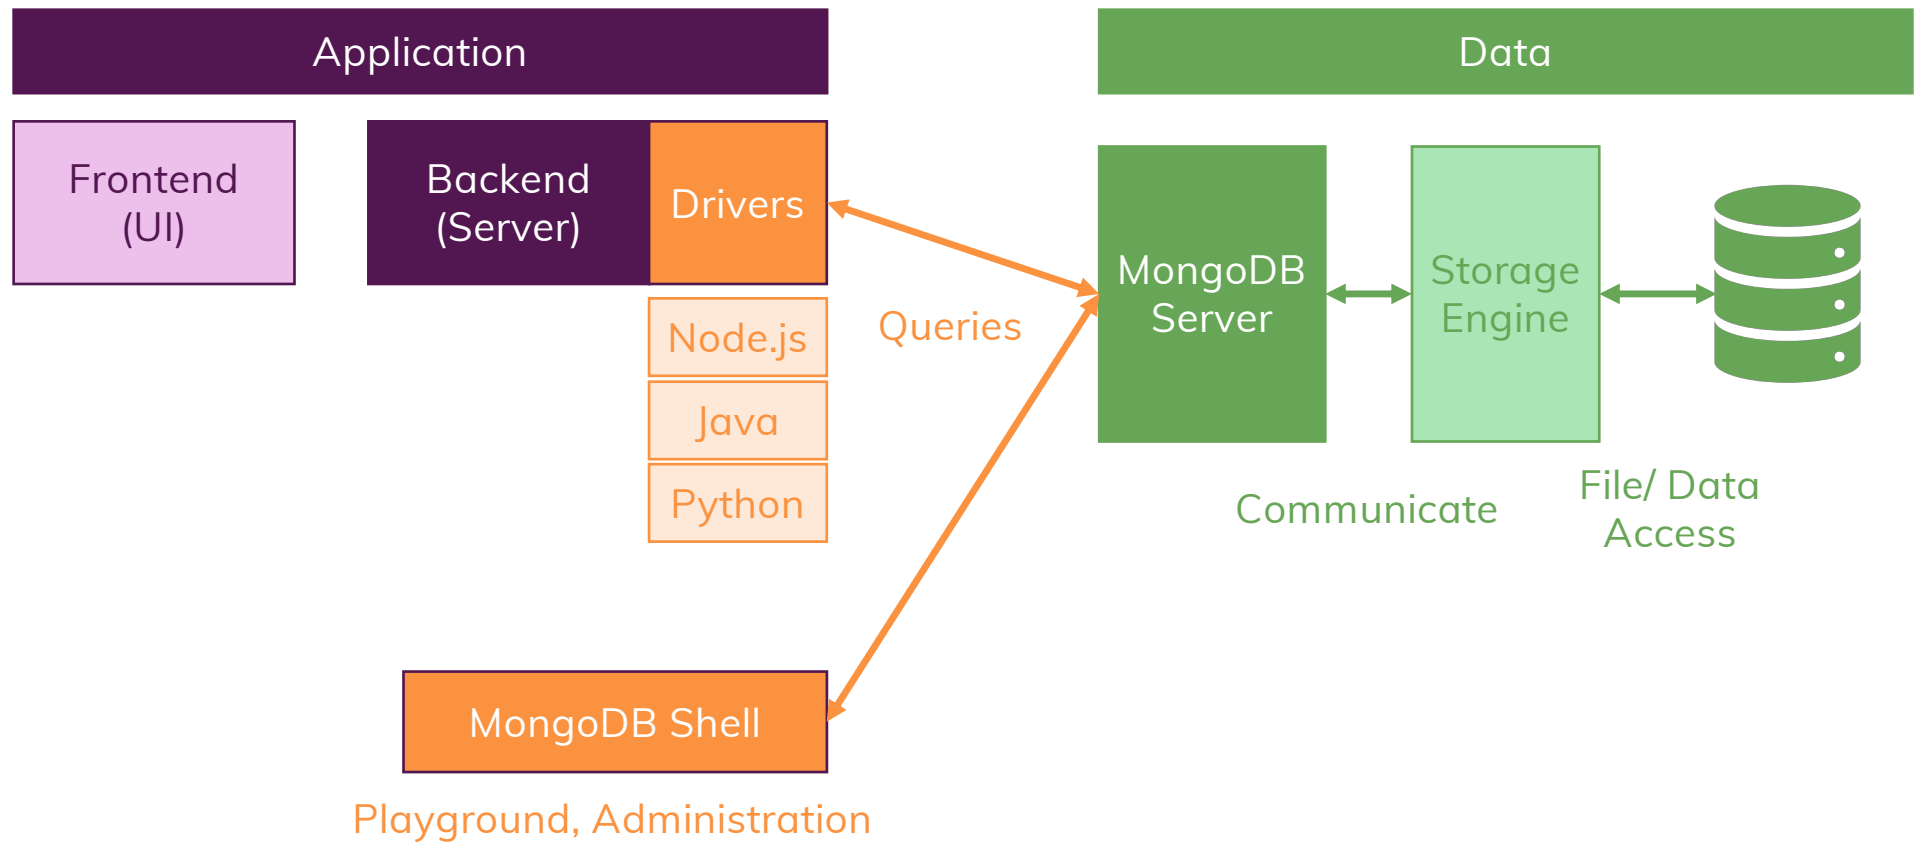
Relational Data needs to be merged manually

Kind of...

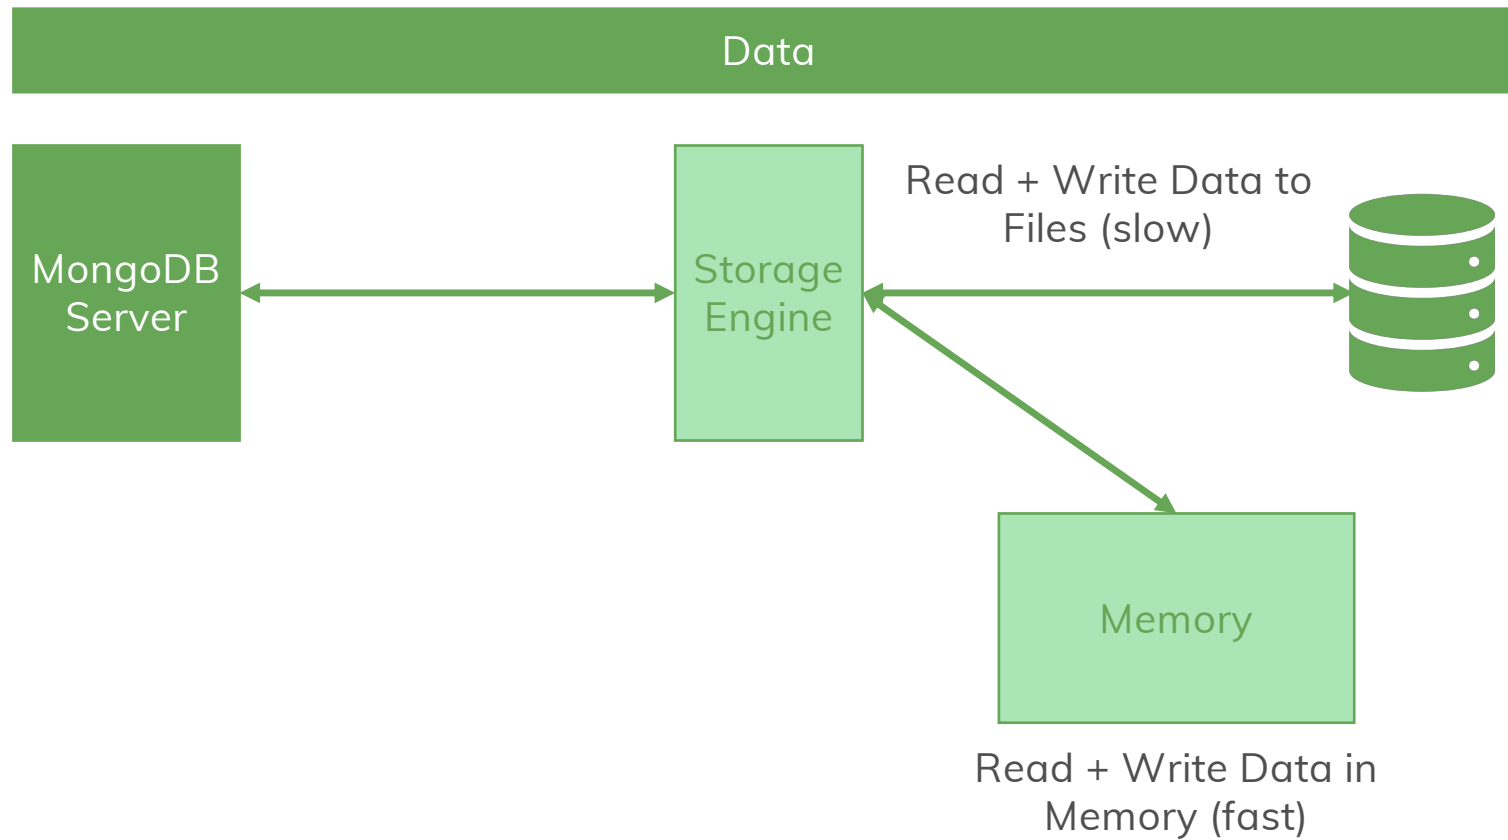
MongoDB Ecosystem



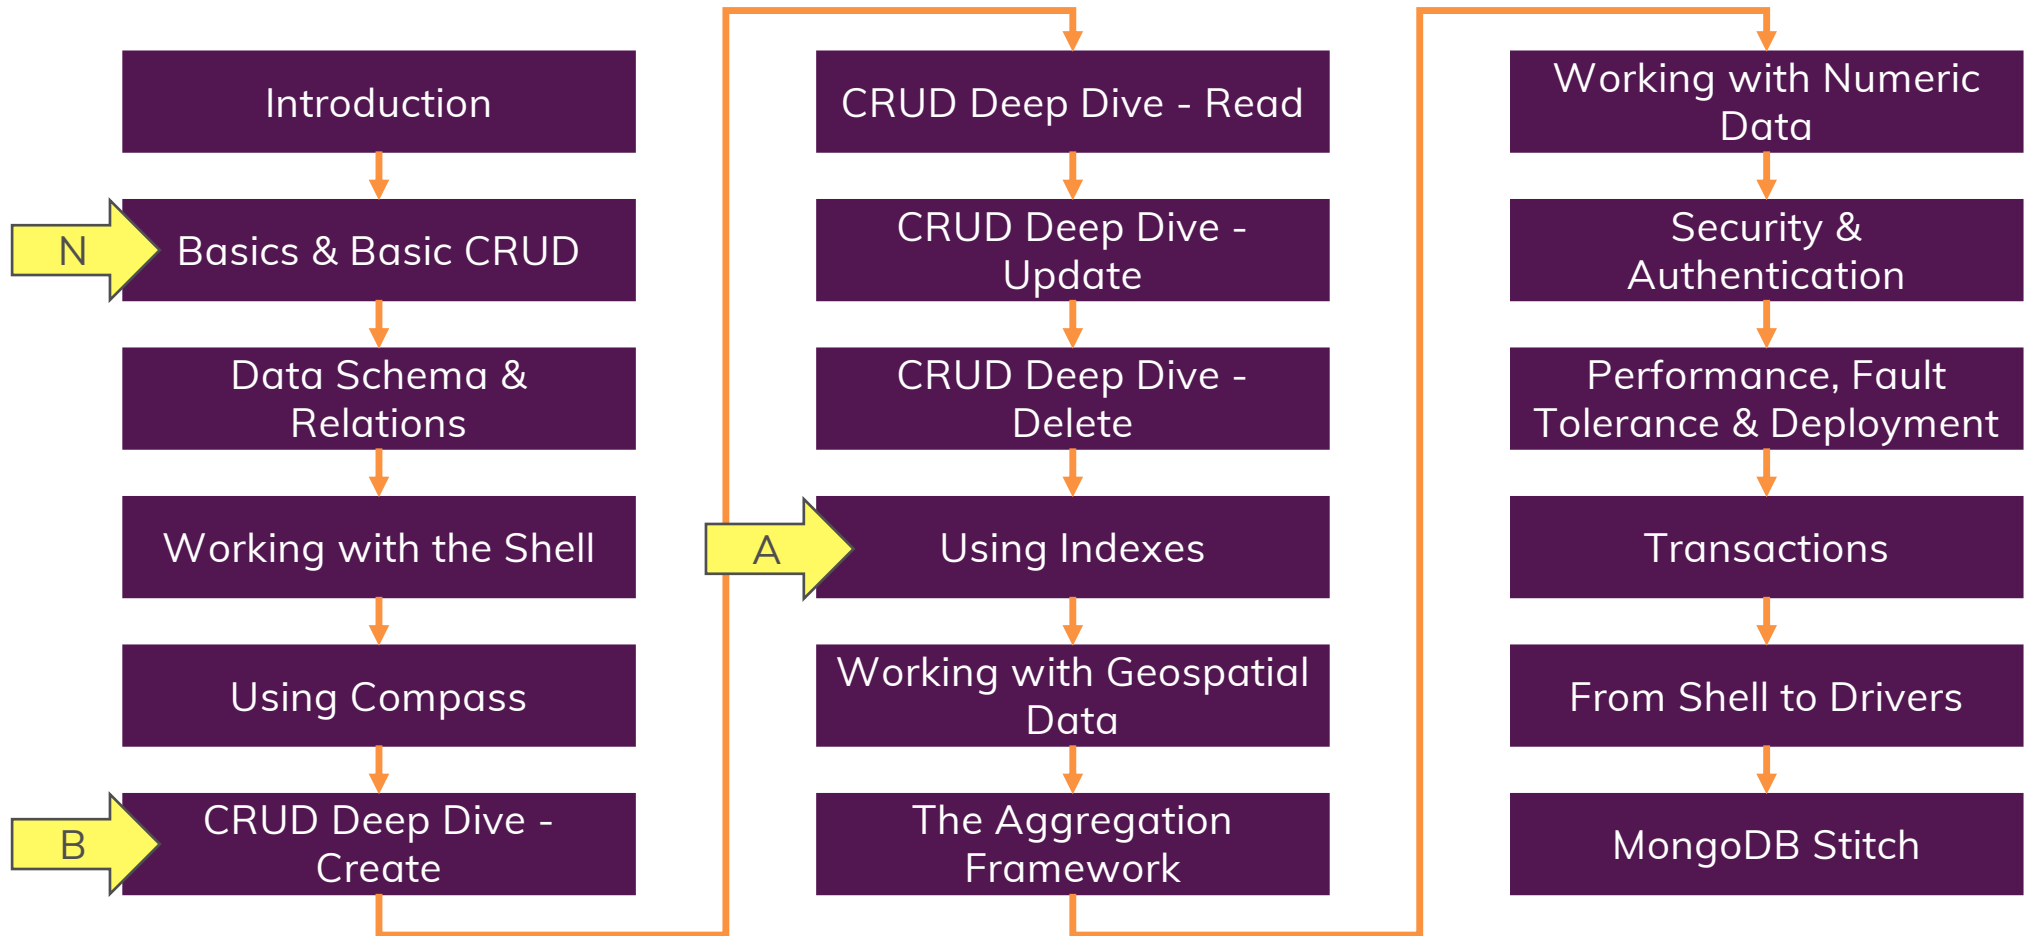
Working with MongoDB



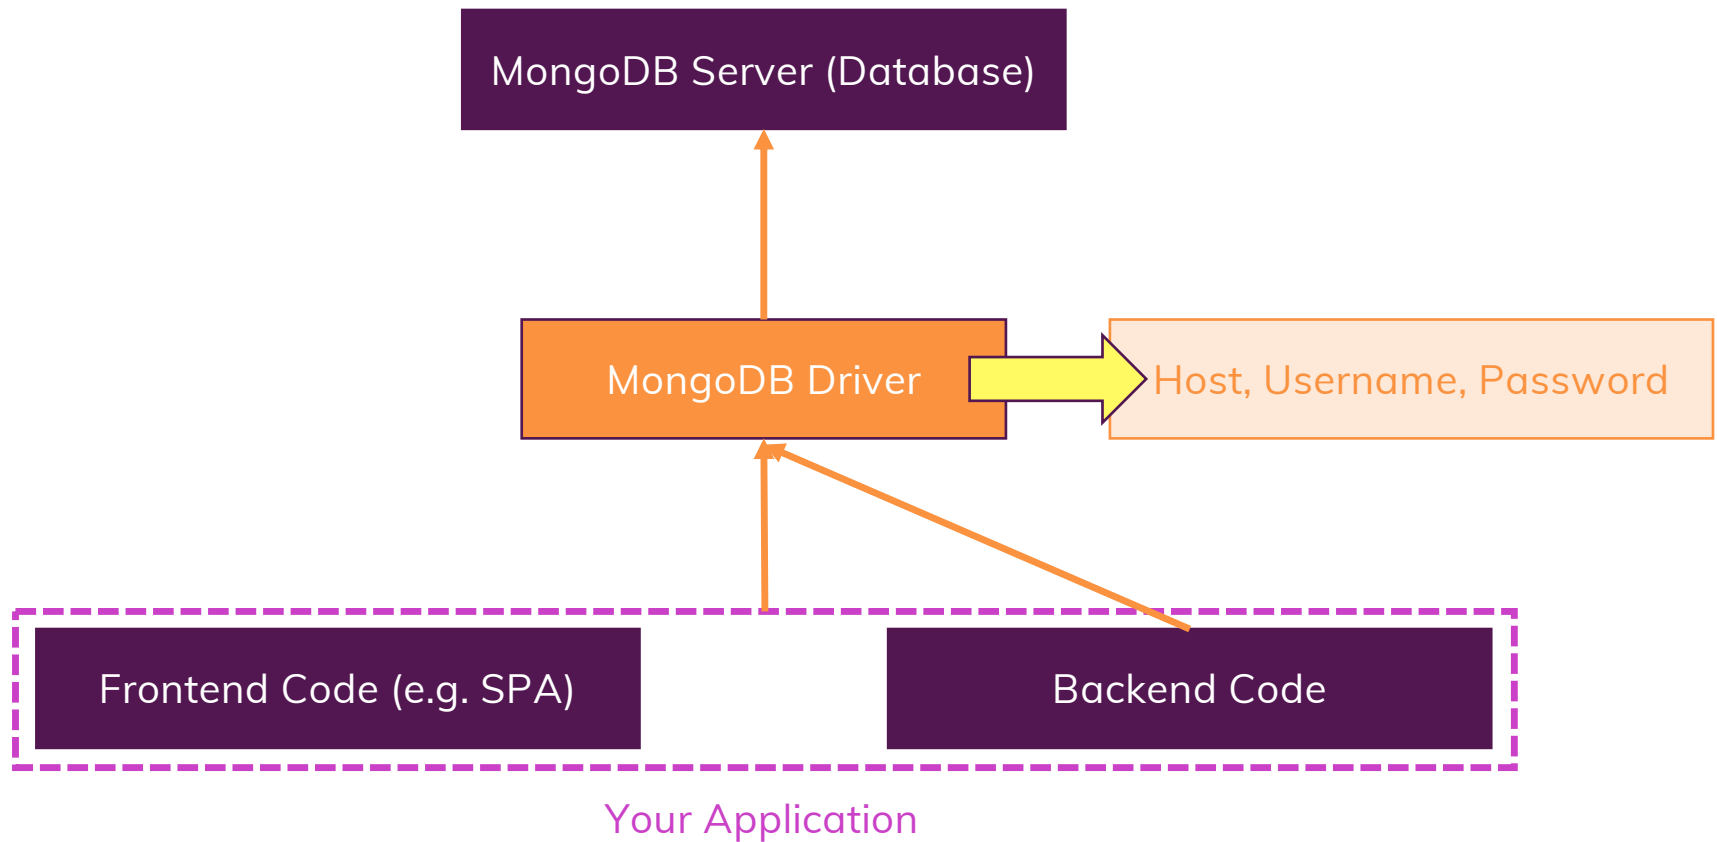
A Closer Look



Outline



Using MongoDB Drivers



How To Get The Most Out Of The Course



Document & CRUD Basics

Working with the Database

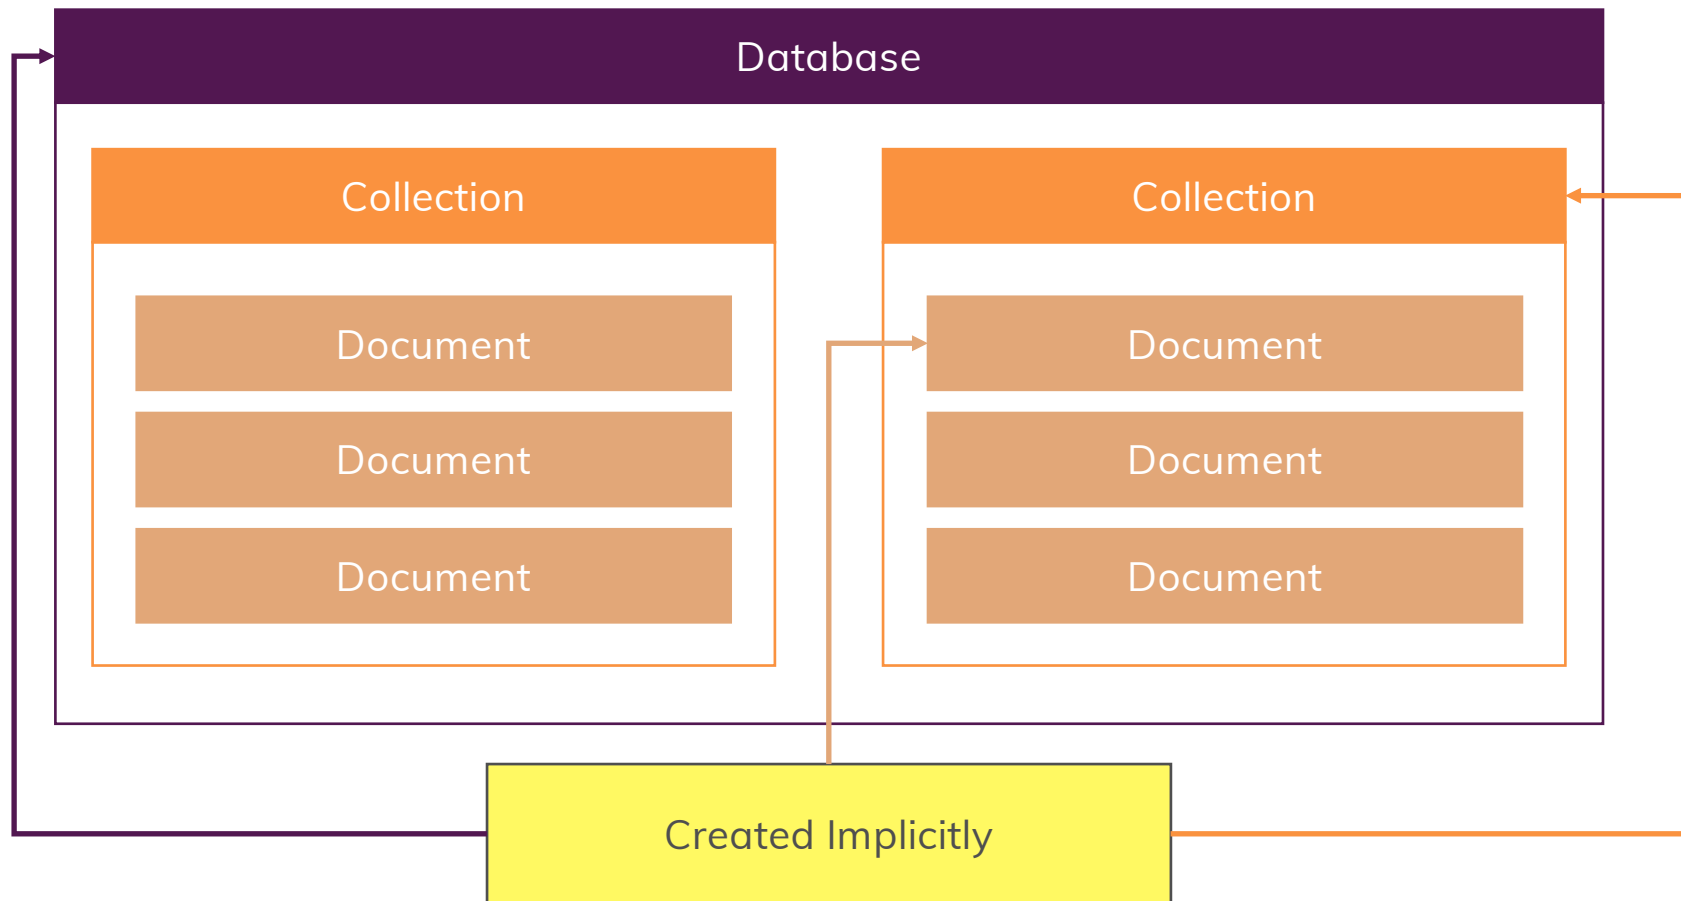
What's Inside This Module?

Basics about Collections & Documents

Basic Data Types

Performing CRUD Operations

Databases, Collections, Documents



JSON

Surrounding curly braces delimit the JSON document

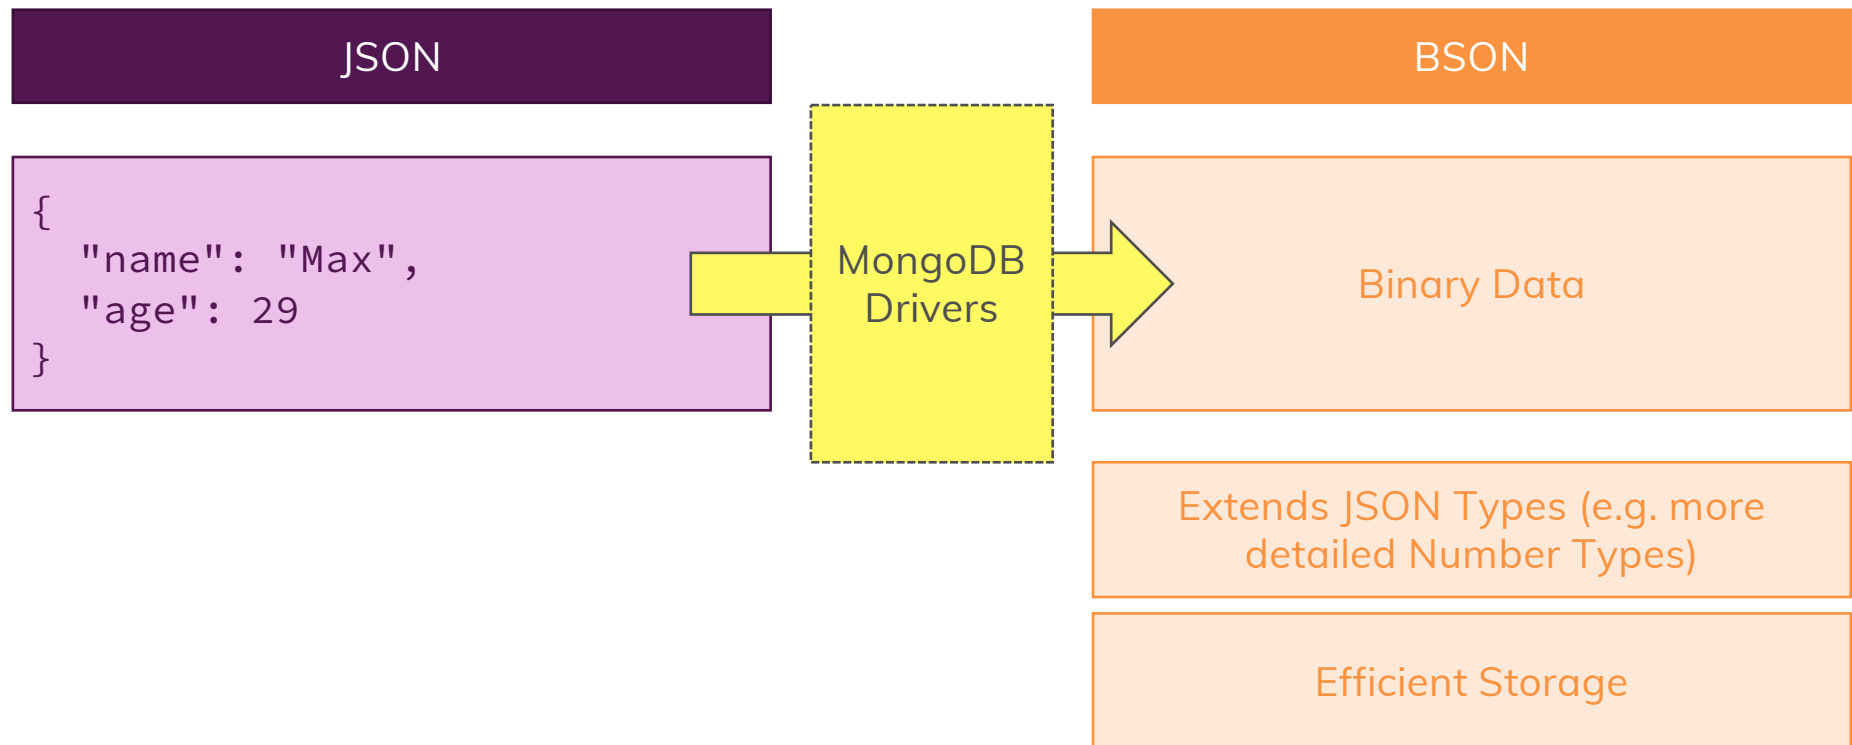
```
{
  "name": "Max",
  "age": 29,
  "isInstructor": true,
  "hobbies": [
    "Sports",
    "Cooking"
  ],
  "address": {
    "street": "My Street 5",
    "city": "Munich"
  }
}
```

This is called a “**Field**” or “**Property**” of the JSON document. Multiple Fields are separated by commas

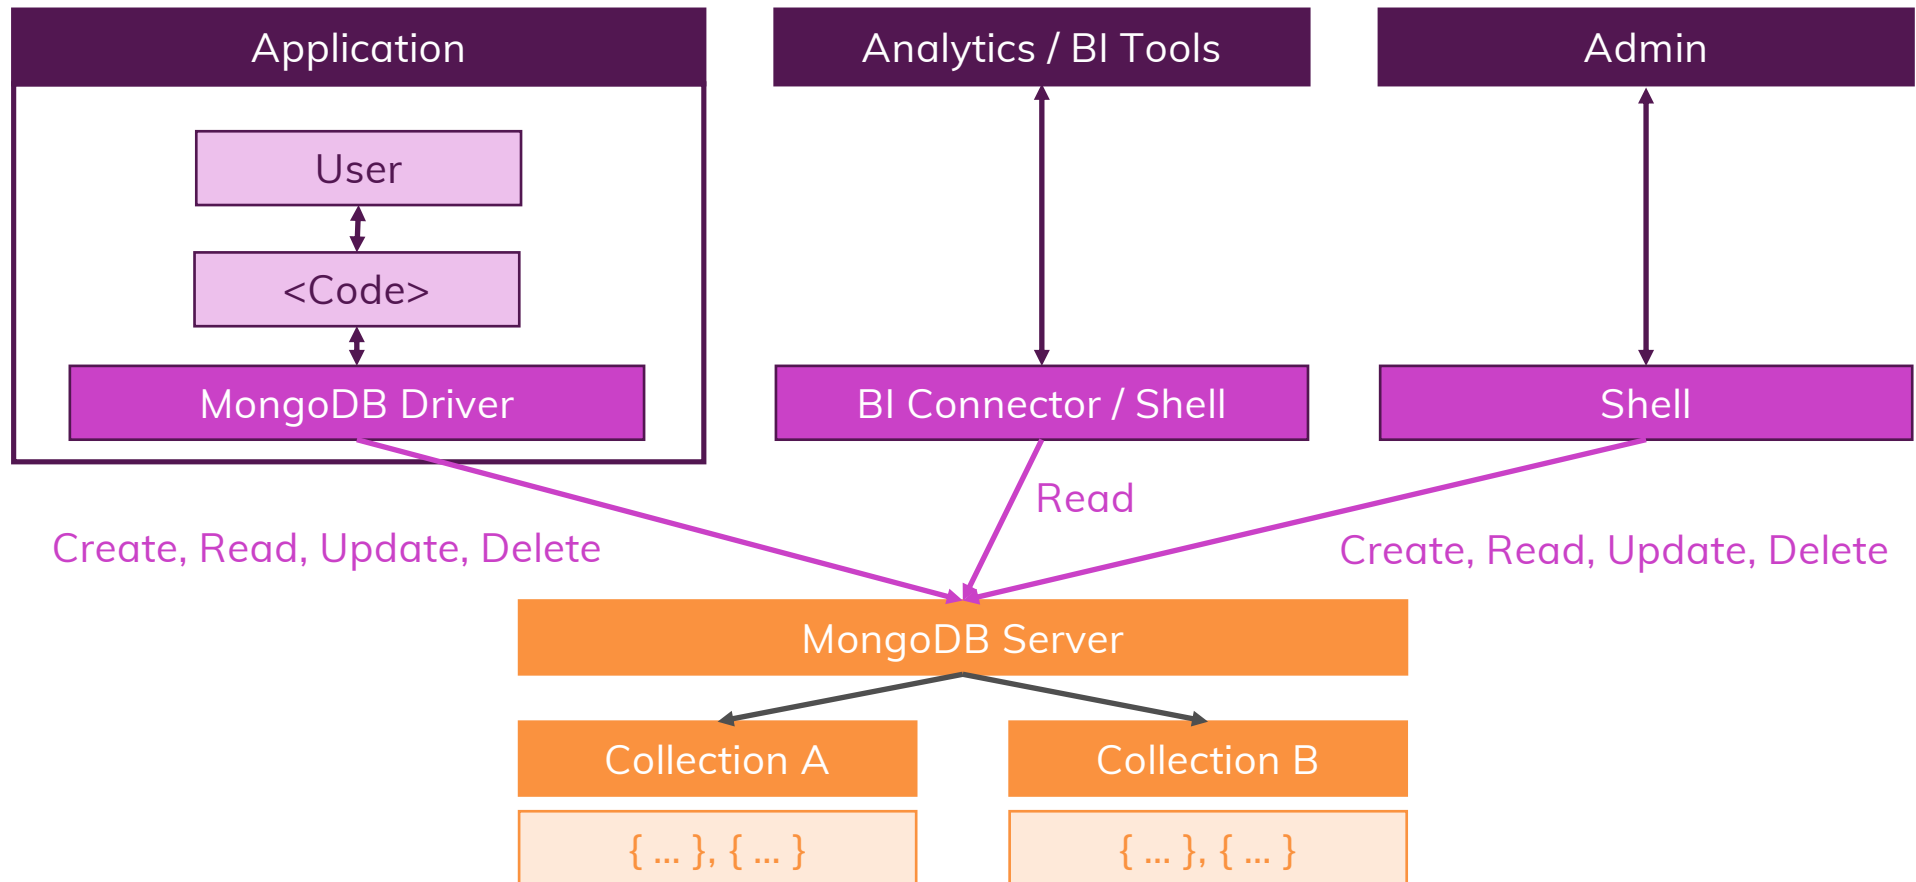
“Fields” consist of a “**Key**” (or “name”) and “**Value**” part. “Key and Value are separated by a colon.

Values can be **strings** (e.g. “Max”), **numbers** (e.g. 29), **booleans** (e.g. true), **arrays** ([...]) and other **documents** (also called objects; { ... })

JSON vs BSON



CRUD Operations & MongoDB



CRUD Operations

Create

`insertOne(data, options)`

`insertMany(data, options)`

Read

`find(filter, options)`

`findOne(filter, options)`

Update

`updateOne(filter, data, options)`

`updateMany(filter, data, options)`

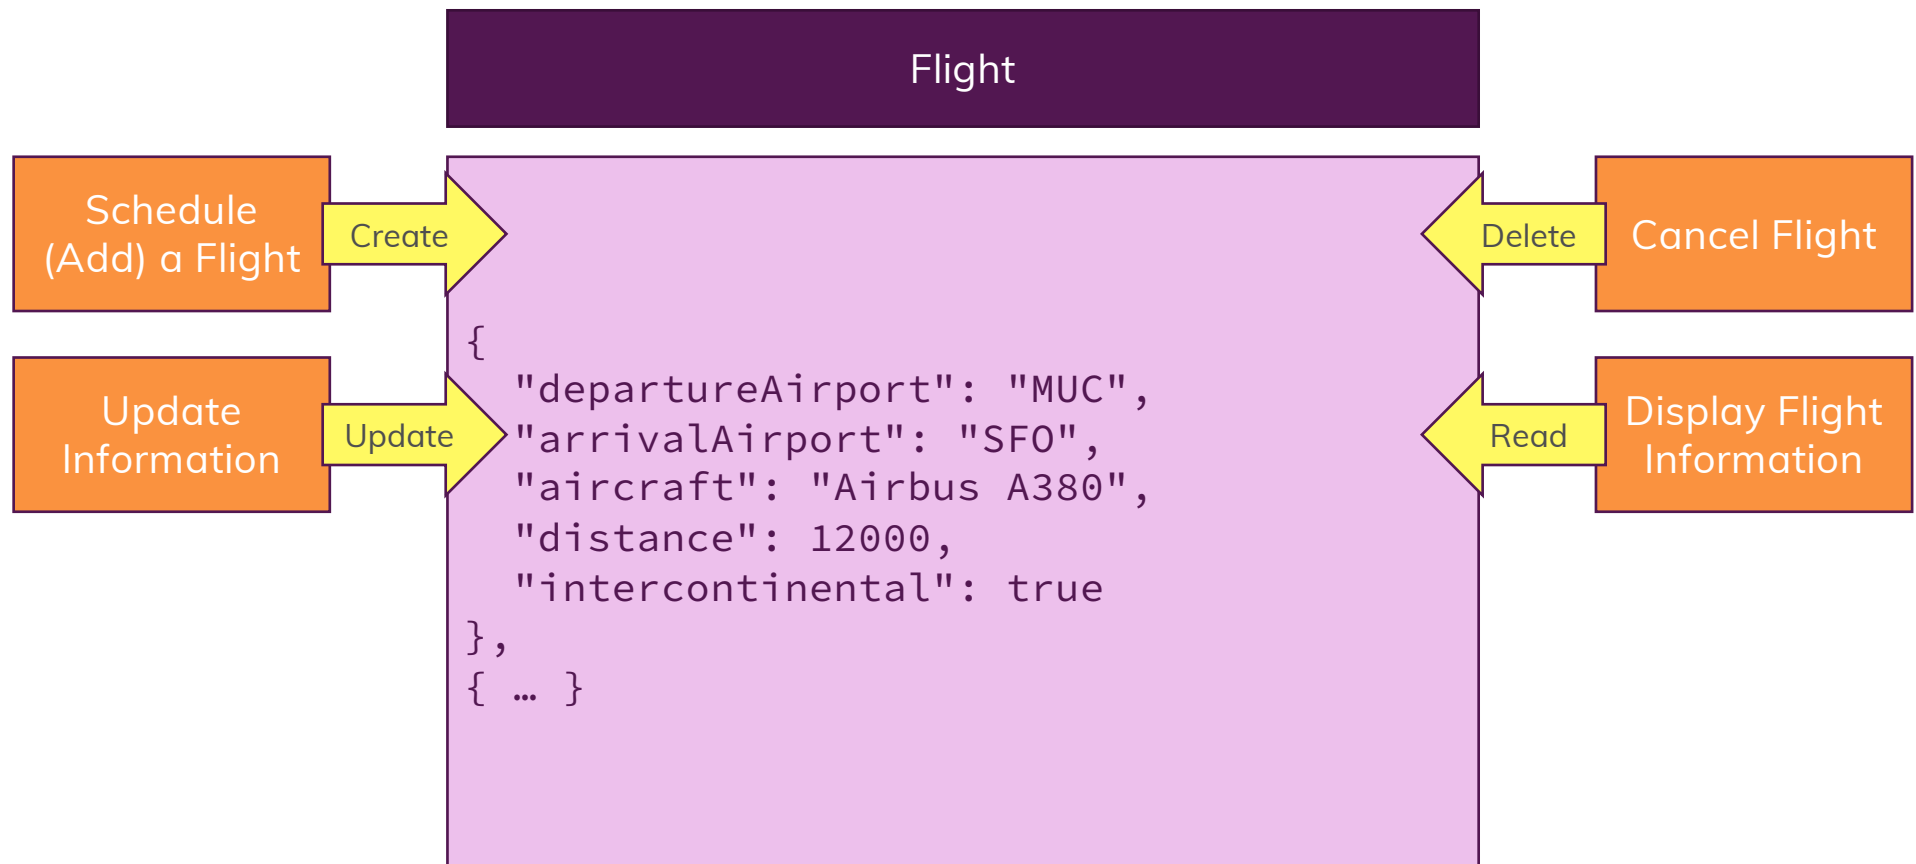
`replaceOne(filter, data, options)`

Delete

`deleteOne(filter, options)`

`deleteMany(filter, options)`

Example #1: Flight Data



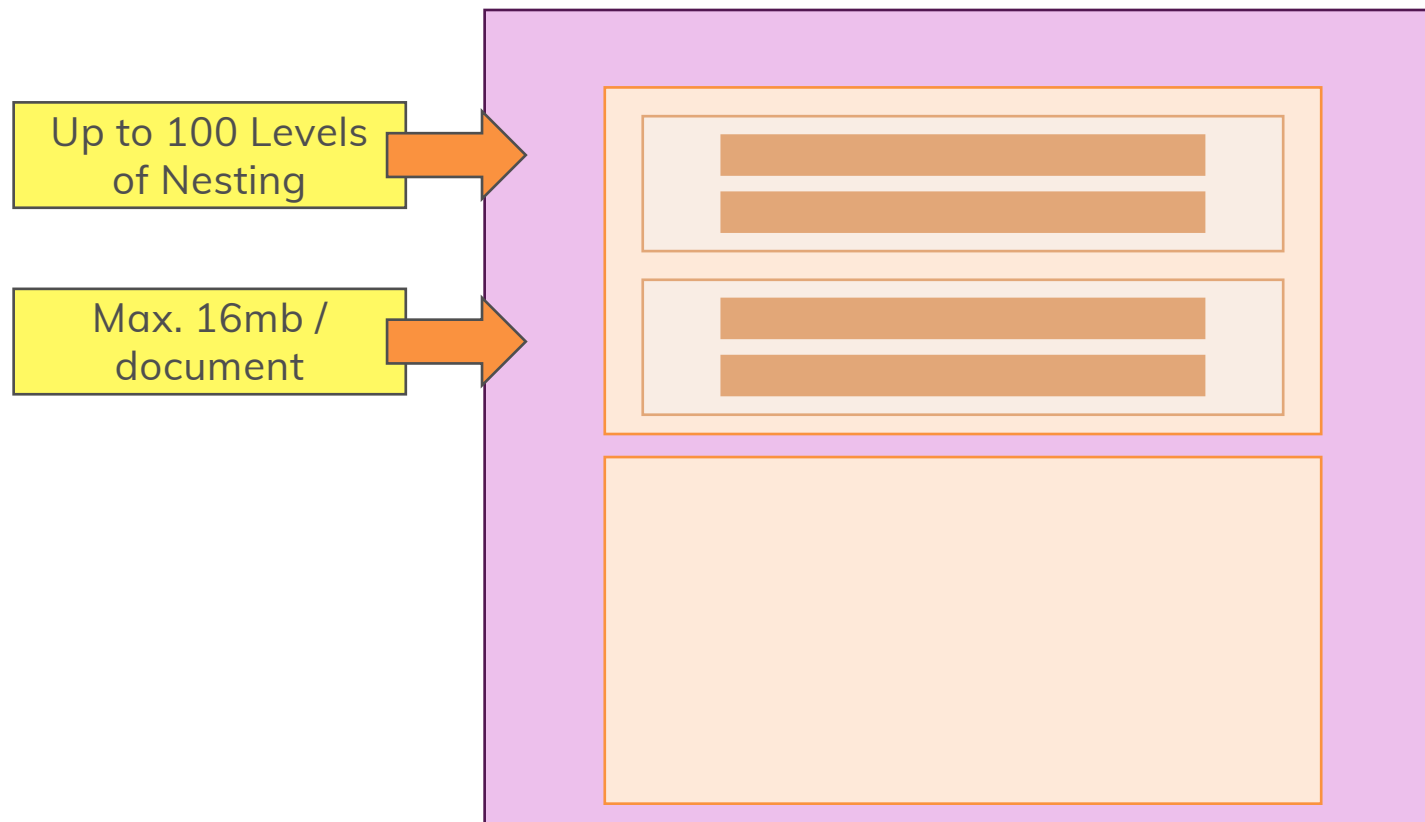
Unique IDs

You MUST have an `_id`


MongoDB creates an `ObjectId()` for you

You can set any other Value


Embedded Documents



Arrays




Cursors




Projection

In Database




```
{  
  "_id": "...",  
  "name": "Max",  
  "age": 29,  
  "job": "instructor"  
}
```





Projection

In Application



```
{  
  "name": "Max",  
  "age": 29  
}
```


update() vs updateOne() vs updateMany()

update()	updateOne()	updateMany()
Overwrite by default	Error without \$set (or other update operators)	Error without \$set (or other update operators)
Use \$set to patch values	Use \$set to patch values	Use \$set to patch values
Update all identified elements	Update first identified element	Update all identified elements
		
	Use these!	

Example #2: Patient Data

Patient

```
{  
  "firstName": "Max",  
  "lastName": "Schwarzmueller",  
  "age": 29,  
  "history": [  
    { "disease": "cold", "treatment": ... },  
    { ... }  
  ]  
}
```

Tasks

1	Insert 3 patient records with at least 1 history entry per patient
2	Update patient data of 1 patient with new age, name and history entry
3	Find all patients who are older than 30 (or a value of your choice)
4	Delete all patients who got a cold as a disease

Module Summary

Databases, Collections, Documents

- A Database holds multiple Collections where each Collection can then hold multiple Documents
- Databases and Collections are created “lazily” (i.e. when a Document is inserted)
- A Document can’t directly be inserted into a Database, you need to use a Collection!

CRUD Operations

- CRUD = Create, Read, Update, Delete
- MongoDB offers multiple CRUD operations for single-document and bulk actions (e.g. `insertOne()`, `insertMany()`, ...)
- Some methods require an argument (e.g. `insertOne()`), others don’t (e.g. `find()`)
- `find()` returns a cursor, NOT a list of documents!
- Use filters to find specific documents

Document Structure

- Each document needs a unique ID (and gets one by default)
- You may have embedded documents and array fields

Retrieving Data

- Use filters and operators (e.g. `$gt`) to limit the number of documents you retrieve
- Use projection to limit the set of fields you retrieve

Data Schemas & Data Modelling

Storing your Data Correctly

What's Inside This Module?

Understanding Document Schemas &
Data Types

Modelling Relations

Schema Validation

Schema-less Or Not?

Isn't MongoDB all about having **NO** data Schemas?



MongoDB enforces no schemas! Documents don't have to use the same schema inside of one collection



But that does not mean that you can't use some kind of schema!

To Schema Or Not To Schema

Chaos!




SQL World!

Products

```
{  
  "title": "Book",  
  "price": 12.99  
}
```

Very Different!

```
{  
  "name": "Bottle",  
  "available": true  
}
```




Products

```
{  
  "title": "Book",  
  "price": 12.99  
}
```

Extra Data

```
{  
  "title": "Bottle",  
  "price": 5.99  
  "available": true  
}
```




Products


```
{  
  "title": "Book",  
  "price": 12.99  
}
```

Full Equality

```
{  
  "title": "Bottle",  
  "price": 5.99  
}
```




Data Types




Data Schemas & Data Modelling

Which Data does my App need or generate?	User Information, Product Information, Orders, ...	Defines the Fields you'll need (and how they relate)
Where do I need my Data?	Welcome Page, Products List Page, Orders Page	Defines your required collections + field groupings
Which kind of Data or Information do I want to display?	Welcome Page: Product Names; Products Page: ...	Defines which queries you'll need
How often do I fetch my data?	For every page reload	Defines whether you should optimize for easy fetching
How often do I write or change my data?	Orders => Often Product Data => Rarely	Defines whether you should optimize for easy writing



Data Schemas & Data Modelling



Relations - Options

Nested / Embedded Documents

Customers

```
{
  userName: 'max',
  age: 29,
  address: {
    street: 'Second Street',
    city: 'New York'
  }
}
```

References

Lots of data duplication!

```
{
  userName: 'max',
  favBooks: [{...}, {...}]
}
```

Customers

```
{
  userName: 'max',
  favBooks: ['id1', 'id2']
}
```

Books

```
{
  _id: 'id1',
  name: 'Lord of the Rings 1'
}
```

Example #1 – Patient <-> Disease Summary



"One patient has one disease summary, a
disease summary belongs to one patient"




Example #2 – Person <-> Car




"One person has one car, a car belongs to one person"



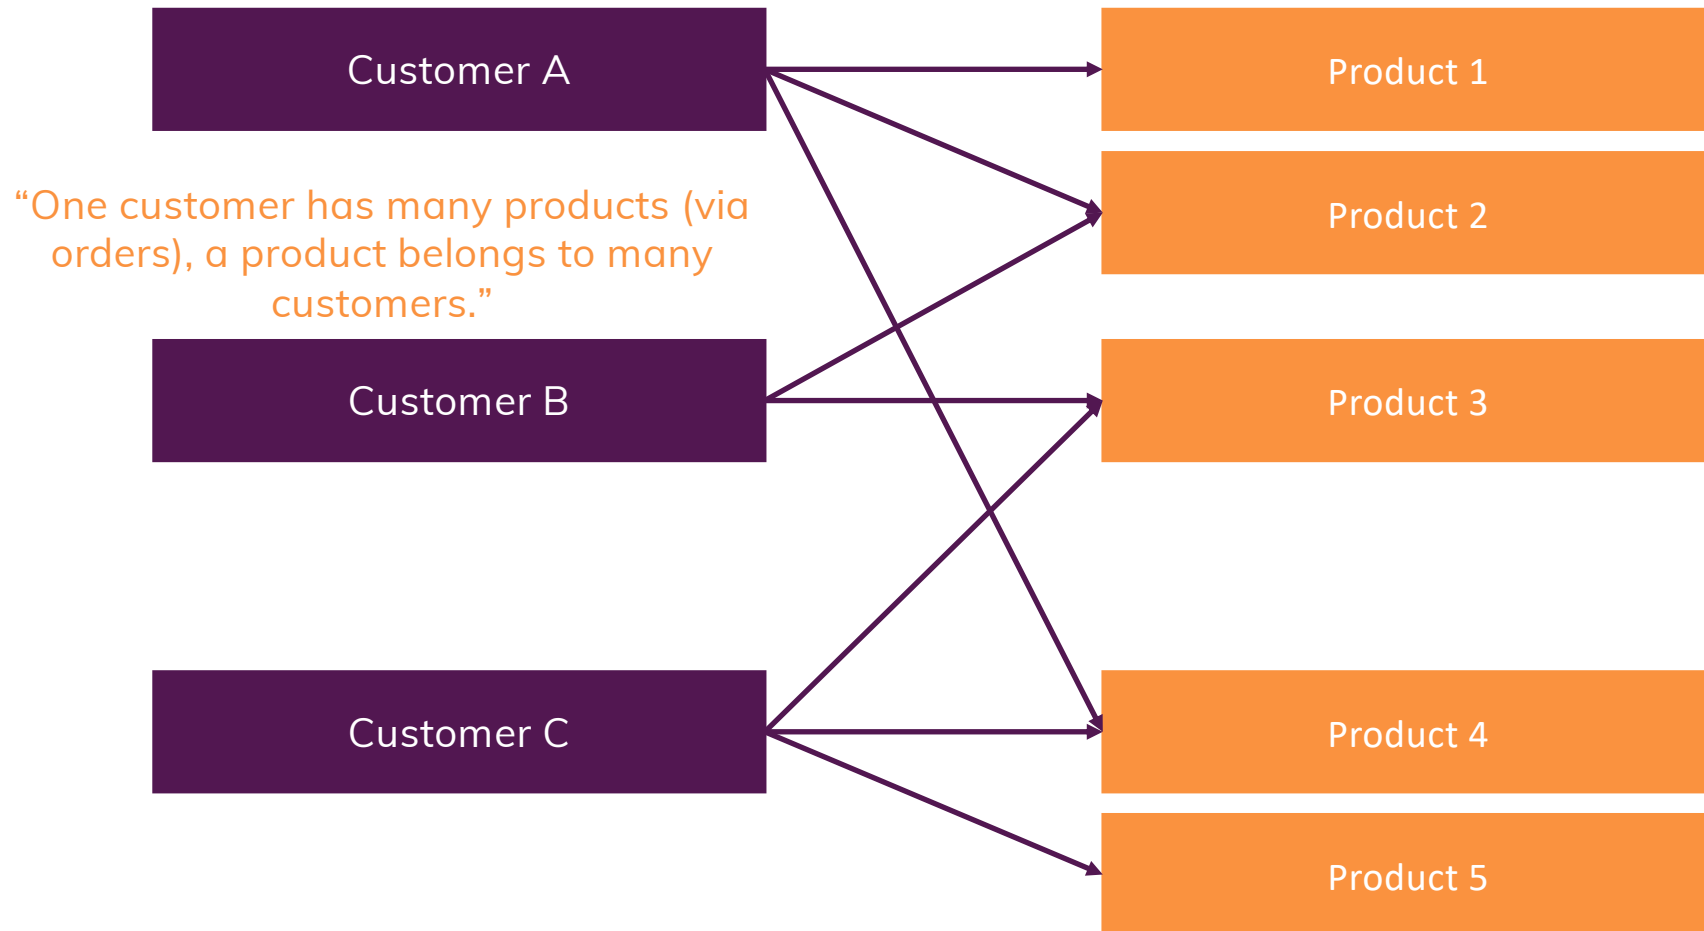
Example #3 – Thread <-> Answers




Example #4 – City <-> Citizens



Example #5 – Customers <-> Products (Orders)



Example #6 – Books <-> Authors



Relations - Options

Nested / Embedded Documents

Group data together logically

Great for data that belongs together and is not really overlapping with other data

Avoid super-deep nesting (100+ levels) or extremely long arrays (16mb size limit per document)


References

Split data across collections

Great for related but shared data as well as for data which is used in relations and standalone


Allows you to overcome nesting and size limits (by creating new documents)

Joining with \$lookup




```
customers.aggregate([
  { $lookup: {
    from: "books",
    localField: "favBooks",
    foreignField: "_id"
    as: "favBookData"
  }
}]
)
```

Example Project: A Blog



Schema Validation



Schema Validation

`validationLevel`

Which documents get validated?

strict



All inserts & updates

moderate



All inserts & updates
to correct documents

`validationAction`

What happens if validation fails?

error



Throw error and deny
insert/ update

warn



Log warning but
proceed

`bypassDocumentValidation()`

Data Modelling & Structuring – Things to Consider

In which Format will you fetch your Data?

How often will you fetch and change your Data?

How much data will you save (and how big is it)?

How is your Data related?

Will Duplicates hurt you (=> many Updates)?

Will you hit Data/ Storage Limits?

Module Summary

Modelling Schemas

- Schemas should be modelled based on your application needs
- Important factors are: Read and write frequency, relations, amount (and size) of data

Schema Validation

- You can define rules to validate inserts and update before writing to the database
- Choose your validation level and action based on your application requirements

Modelling Relations

- Two options: Embedded documents or references
- Use embedded documents if you got one-to-one or one-to-many relationships and no app or data size reason to split
- Use references if data amount/ size or application needs require it or for many-to-many relations
- Exceptions are always possible => Keep your app requirements in mind!

Working with Shell & Server

Beyond Start & Stop

What's Inside This Module?

Start MongoDB Server as Process &
Service

Configuring Database & Log Path (and
Mode)

Fixing Issues

Diving Deeper Into **CREATE**

A Closer Look at Creating & Importing Documents

What's Inside This Module?

Document Creation Methods (CREATE)

Importing Documents

CREATE Documents

`insertOne()`

```
db.collectionName.insertOne({field: "value"})
```

`insertMany()`

```
db.collectionName.insertMany([  
    {field: "value"},  
    {field: "value"}])
```


`insert()`

```
db.collectionName.insert()
```

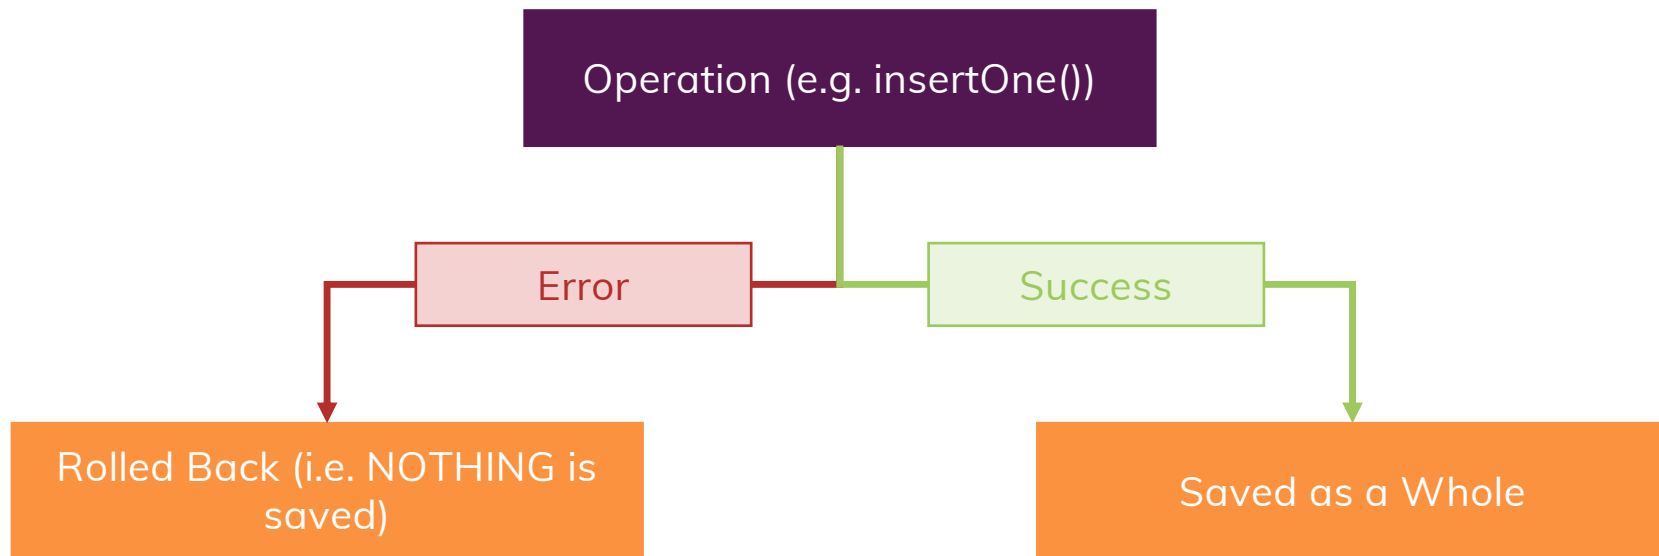
`mongoimport`

```
mongoimport -d cars -c carsList --drop --jsonArray
```

WriteConcern



What is “Atomicity”?



MongoDB CRUD Operations are Atomic on the Document Level (including Embedded Documents)

Tasks

1

Insert multiple companies (company data of your choice) into a collection – both with insertOne() and insertMany()

2

Deliberately insert duplicate ID data and “fix” failing additions with unordered inserts

3

Write data for a new company with both journaling being guaranteed and not being guaranteed

Module Summary

insertOne(), insertMany()

- You can insert documents with insertOne() (one document at a time) or insertMany() (multiple documents)
- insert() also exists but it's not recommended to use it anymore – it also doesn't return the inserted ids

WriteConcern

- Data should be stored and you can control the “level of guarantee” of that to happen with the writeConcern option
- Choose the option value based on your app requirements

Ordered Insertes

- By default, when using insertMany(), inserts are ordered – that means, that the inserting process stops if an error occurs
- You can change this by switching to “unordered inserts” – your inserting process will then continue, even if errors occurred
- In both cases, no successful inserts (before the error) will be rolled back

READing Documents with Operators

Accessing the Required Data Efficiently


What's Inside This Module?

Methods, Filters & Operators

Query Selectors (READ)

Projection Operators (READ)

Methods, Filters & Operators



Operators

Read

Update

Query & Projection

Update

Query Modifiers

Aggregation

Query Selectors

Fields

Change Stream
Operators




Pipeline Stages

Projection Operators


Arrays

Aggregation
Module
Operators

How Operators Impact our Data

Type	Purpose	Changes Data?	Example
Query Operator	Locate Data		\$eq
Projection Operator	Modify data presentation		\$
Update Operator	Modify + add additional data		\$inc

Query Selectors & Projection Operators



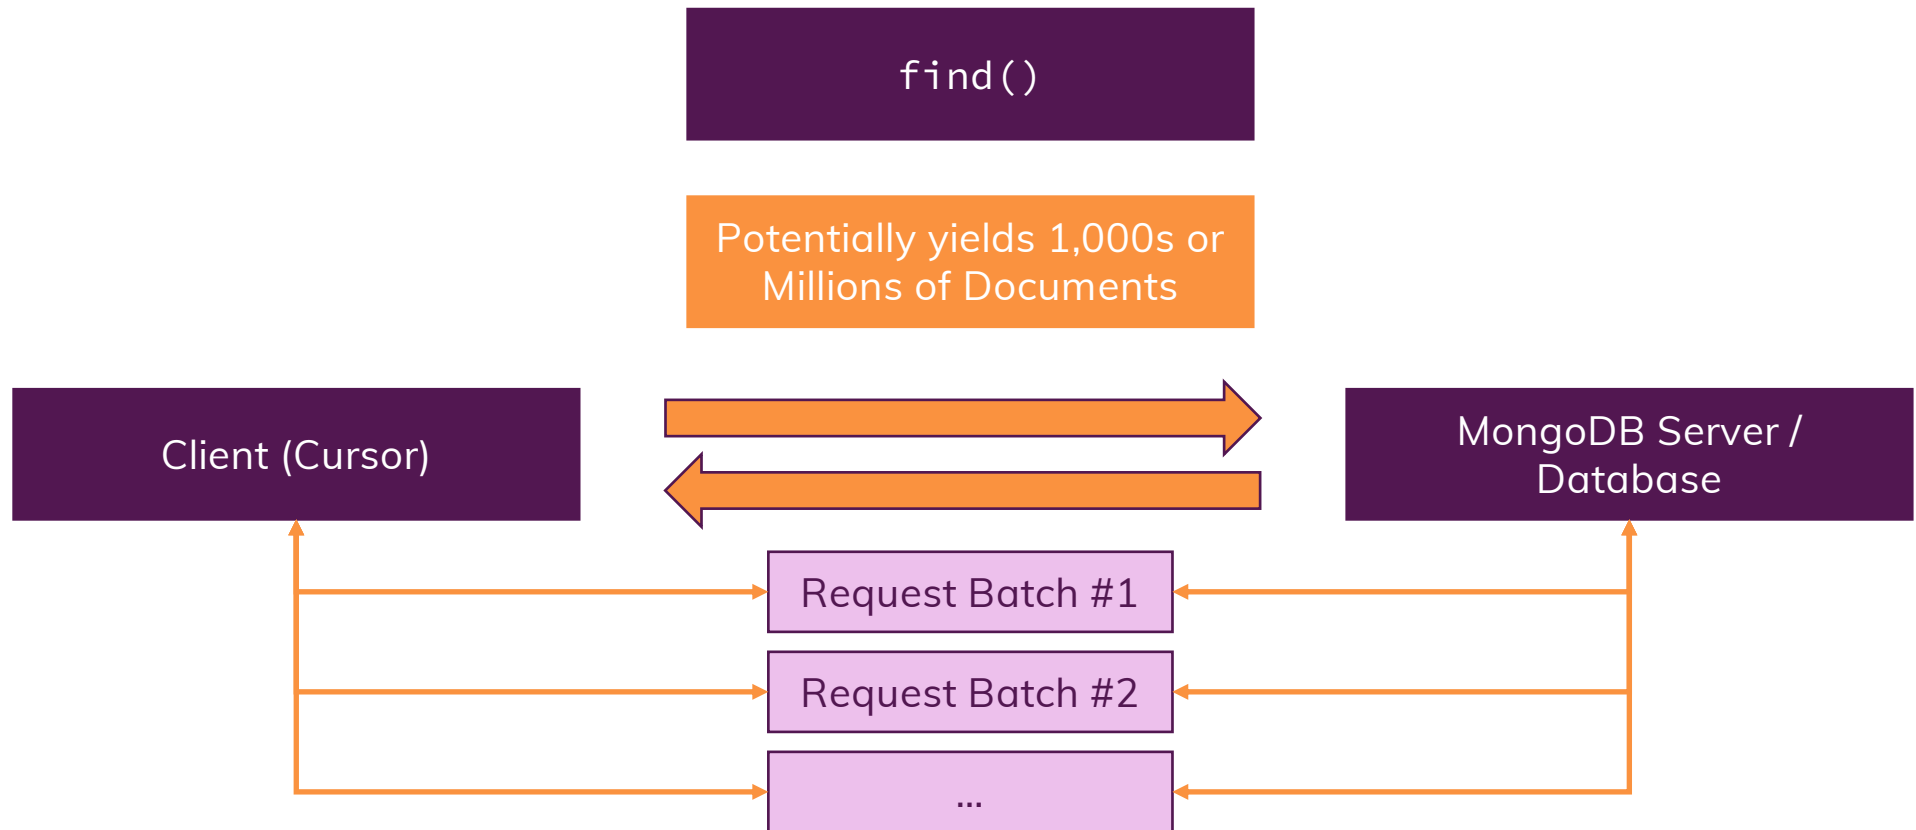
Tasks

- | | |
|---|---|
| 1 | Import the attached data into a new database (e.g. boxOffice) and collection (e.g. movieStarts) |
| 2 | Search all movies that have a rating higher than 9.2 and a runtime lower than 100 minutes |
| 3 | Search all movies that have a genre of “drama” or “action” |
| 4 | Search all movies where visitors exceeded expectedVisitors |

Tasks

- | | |
|---|--|
| 1 | Import the attached data file into a new collection (e.g. exmoviestarts) in the boxOffice database |
| 2 | Find all movies with exactly two genres |
| 3 | Find all movies which aired in 2018 |
| 4 | Find all movies which have ratings greater than 8 but lower than 10 |

Understanding Cursors



Tasks

- 1 For this assignment, we'll work on the "extended boxoffice" dataset (which was imported in the previous assignment)
- 2 Filter for any data of your choice (e.g. all data) and make sure to only include title + visitors in your result data.
- 3 Search for all movies that have an entry of 10 in their ratings array and return just that array entry (inside of the array) in the result data
- 4 Repeat step 3) but return all "action" genre entries instead

Module Summary

Query Selectors & Operators

- You can read documents with `find()` and `findOne()`
- `find()` returns a cursor which allows you to fetch data step-by-step
- Both `find()` and `findOne()` take a filter (optional) to narrow down the set of documents they return
- Filters can use a variety of query selectors/ operators to control which documents are retrieved

Cursors

- `find()` returns a cursor to allow you to efficiently retrieve data step by step (instead of fetching all the documents in one step)
- You can use a cursor to move through the documents
- `sort()`, `skip()` and `limit()` can be used to control the order, portion and quantity of the retrieved results

Projection

- Projection allows you to control which fields are returned in your result set
- You can include fields (`field: 1`) and exclude them (`field: 0`)
- For arrays, special projection operators help you return the right field data

Understanding Document **U**PDATES

Because we Always need the Latest Information

What's Inside This Module?

Document Updating Operator (**U**PDATE)

Updating Fields

Updating Arrays

Operators

Read

Update

Query & Projection

Update

Query Modifiers

Aggregation

Query Selectors

Fields

Change Stream
Deprecated

Pipeline Stage




Projection Operators

Arrays

Pipeline Module
Pipeline Operators

Bitwise

How Operators Impact our Data

Type	Purpose	Changes Data?	Example
Query Operator	Locate Data		\$eq
Projection Operator	Modify data presentation		\$elemMatch
Update Operator	Modify + add additional data		\$rename

Update Operators

Operators		Operator Examples	
Fields		<code>\$currentDate</code>	<code>\$mul</code>
Arrays	Operators	<code>\$push</code>	<code>\$pop</code>
	Modifiers	<code>\$position</code>	<code>\$slice</code>

Tasks

1

Create a new collection ("sports") and upsert two new documents into it (with these fields: "title", "requiresTeam")

2

Update all documents which do require a team by adding a new field with the minimum amount of players required

3

Update all documents that require a team by increasing the number of required players by 10

Module Summary

updateOne() & updateMany()

- You can use updateOne() and updateMany() to update one or more documents in a collection
- You specify a filter (query selector) with the same operators you know from find()
- The second argument then describes the update (e.g. via \$set or other update operators)

Update Operators

- You can update fields with a broad variety of field update operators like \$set, \$inc, \$min etc
- If you need to work on arrays, take advantage of the shortcuts (\$, \$[] and \$[<identifier>] + arrayFilters)
- Also use array update operators like \$push or \$pop to efficiently add or remove elements to or from arrays

Replacing Documents

- Even though it was not covered again, you also learned about replaceOne() earlier in the course – you can use that if you need to entirely replace a doc

DELETE Documents

Sometimes we have to Get Rid of Data

What's Inside This Module?

Document Deletion Methods (DELETE)

Indexes

Retrieving Data Efficiently


What's Inside This Module?

What are Indexes?

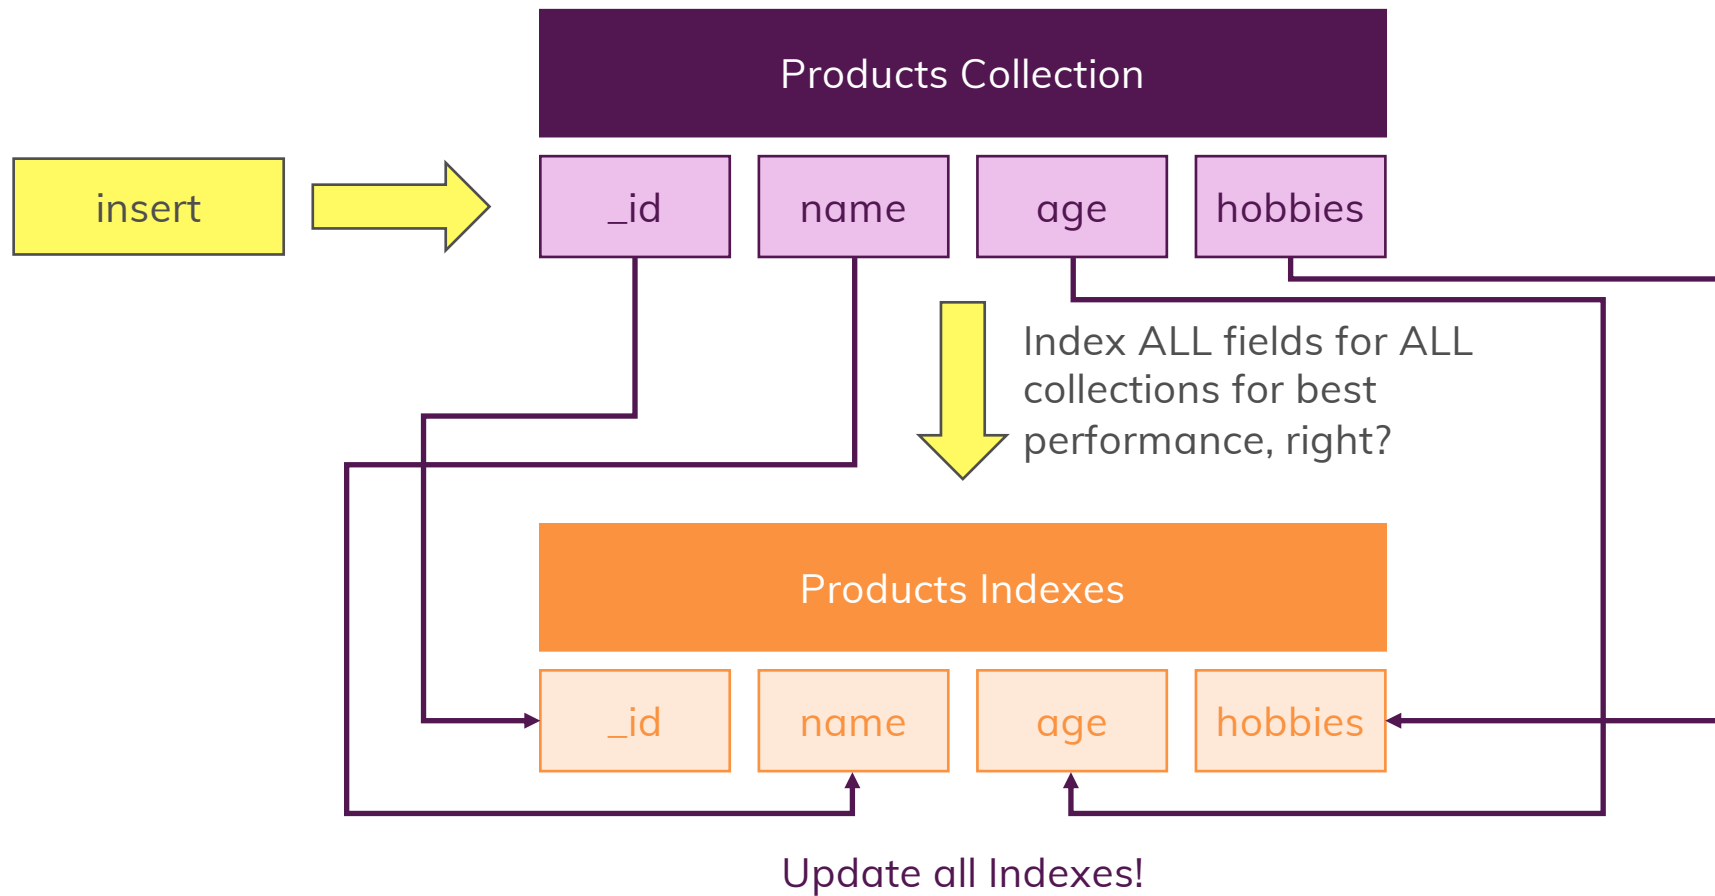
Different Types of Indexes

Using & Optimizing Indexes

Why Indexes?



Don't Use Too Many Indexes!



Index Types

“Normal”

Ordered field

{ name: 1 }

Compound

Multiple, combined ordered fields

{ name: 1, age: -1 }

Multikey

Ordered array values

{ hobbies: 1 }

Text

Ordered text fragments

{ description: "text" }

Geospatial

Ordered geodata

{ location: "2d" }

Index Config

Custom Name


Unique

Partial


Sparse

TTL


Query Diagnosis & Query Planning



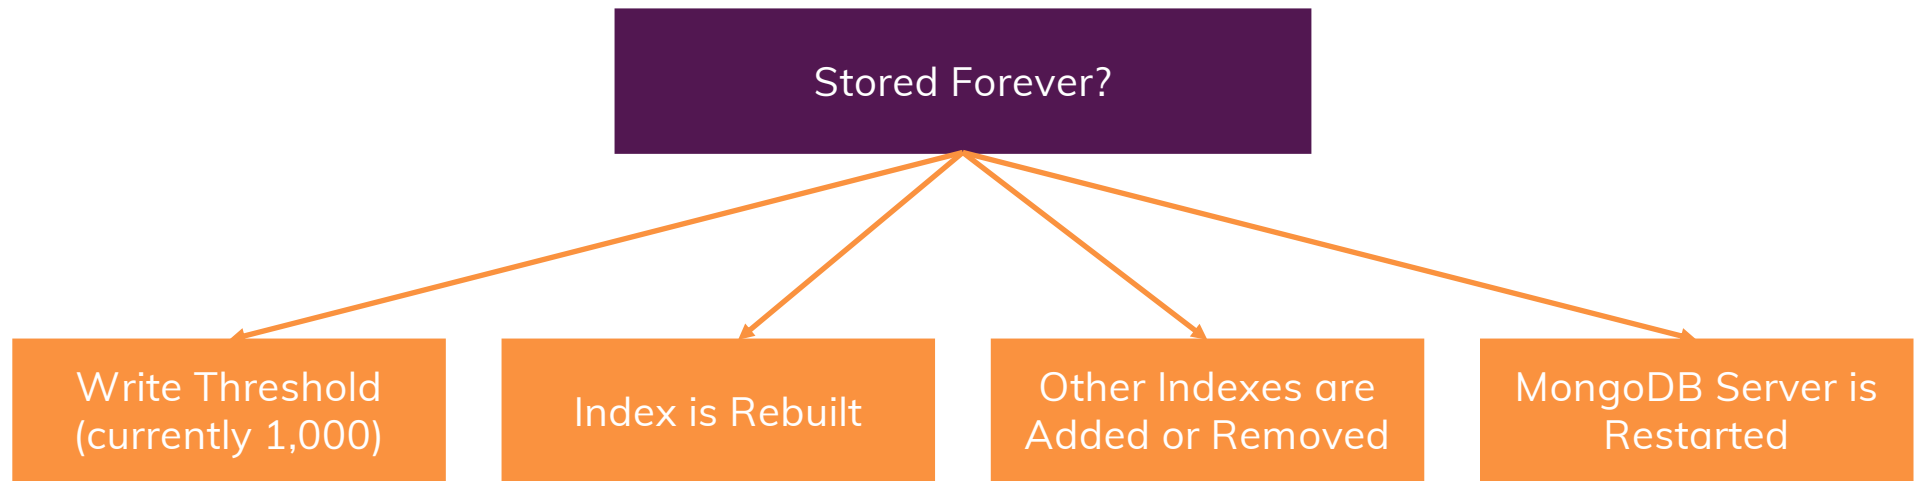
Efficient Queries & Covered Queries



“Winning Plans”



Clearing the Winning Plan from Cache



Understanding "text" Indexes

This product is a must-buy for all fans of modern fiction!



Text Index

product

must

buy

fans

modern

fiction

Stopwords (e.g. "a") are eliminated!

Building Indexes

Foreground

Collection is locked during
index creation

Faster

Background

Collection is accessible during
index creation

Slower

Module Summary

What and Why?

- Indexes allow you to retrieve data more efficiently (if used correctly) because your queries only have to look at a subset of all documents
- You can use single-field, compound, multi-key (array) and text indexes
- Indexes don't come for free, they will slow down your writes

Queries & Sorting

- Indexes can be used for both queries and efficient sorting
- Compound indexes can be used as a whole or in a "left-to-right" (prefix) manner (e.g. only consider the "name" of the "name-age" compound index)

Query Diagnosis & Planning

- Use `explain()` to understand how MongoDB will execute your queries
- This allows you to optimize both your queries and indexes

Index Options

- You can also create TTL, unique or partial indexes
- For text indexes, weights and a `default_language` can be assigned

Geospatial Queries

Finding Places

What's Inside This Module?

Storing Geospatial Data in GeoJSON
Format

Querying Geospatial Data

Tasks

1	Pick 3 Points on Google Maps and store them in a collection
2	Pick a point and find the nearest points within a min and max distance
3	Pick an area and see which points (that are stored in your collection) it contains
4	Store at least one area in a different collection
5	Pick a point and find out which areas in your collection contain that point

Module Summary

Storing Geospatial Data

- You store geospatial data next to your other data in your documents
- Geospatial data has to follow the special GeoJSON format – and respect the types supported by MongoDB
- Don't forget that the coordinates are [longitude, latitude], not the other way around!

Geospatial Queries

- \$near, \$geoWithin and \$geoIntersects get you very far
- Geospatial queries work with GeoJSON data

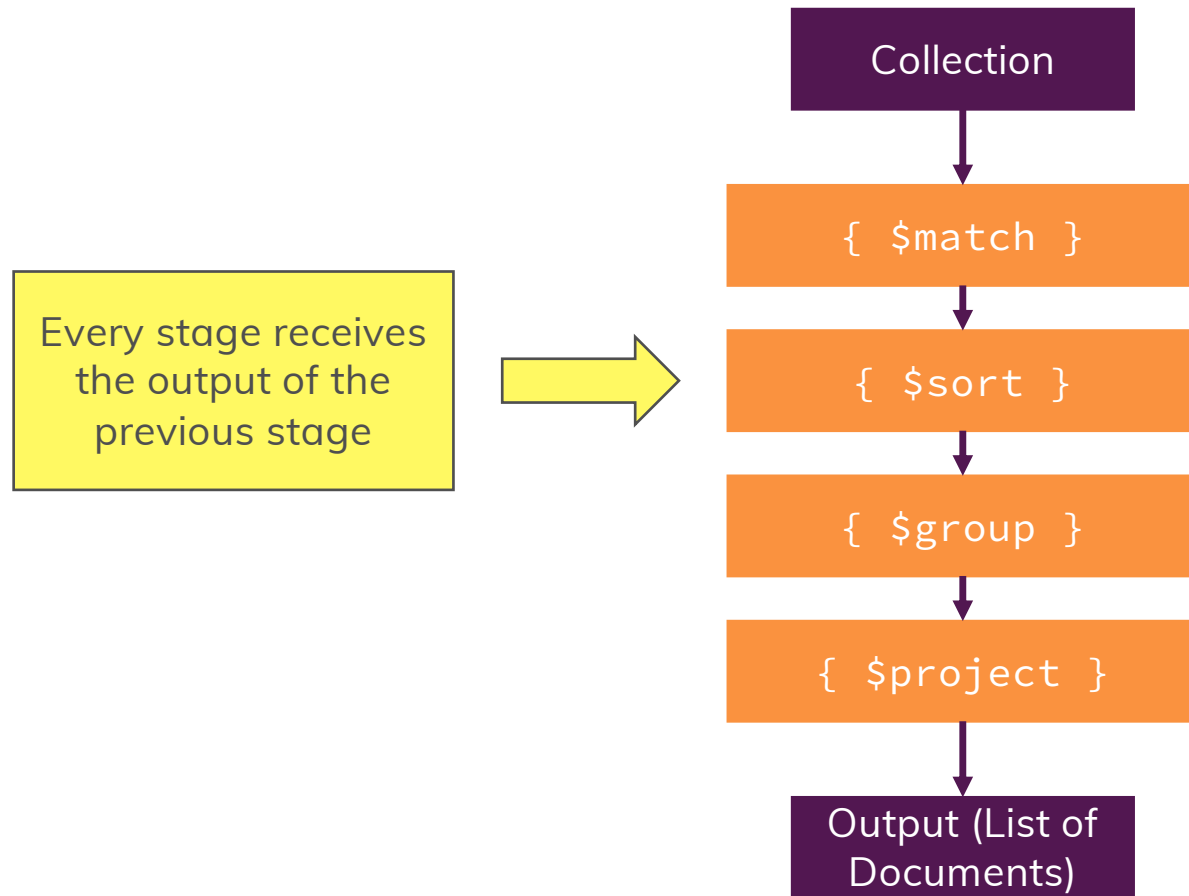
Geospatial Indexes

- You can add an index to geospatial data: "2dsphere"
- Some operations (\$near) require such an index

Using the Aggregation Framework

Retrieving Data Efficiently & In a Structured Way

What is the “Aggregation Framework”?




Pipeline Stages

[Check official docs](#)

\$group vs \$project

\$group

n:1



Sum, Count, Average, Build Array

\$project

1:1



Include/ Exclude Fields, Transform Fields (within a Single Document)

\$unwind

```
{ name: "Max", hobbies: ["Sports", "Cooking"] }
```



```
{ name: "Max", hobbies: "Sports" }
```

```
{ name: "Max", hobbies: "Cooking" }
```

\$skip + \$limit + \$sort

The Order Matters!



\$text

Do a Text Index Search

Has to be the First Pipeline
Stage!

Aggregation Pipeline Optimization

MongoDB automatically optimizes for you!

Module Summary

Stages & Operators

- There are plenty of available stages and operators you can choose from
- Stages define the different steps your data is funneled through
- Each stage receives the output of the last stage as input
- Operators can be used inside of stages to transform, limit or re-calculate data

Important Stages

- The most important stages are \$match, \$group, \$project, \$sort and \$unwind – you'll work with these a lot
- Whilst there are some common behaviors between find() filters + projection and \$match + \$project, the aggregation stages generally are more flexible

Working with Numeric Data

More Complex Than You Might Think

Integers, Longs, Doubles

Integers (int32)

Only full Numbers

-2,147,483,648
to
2,147,483,647

Use for “normal”
integers

Longs (int64)

Only full Numbers

-9,223,372,036,854,
775,808
to
9,223,372,036,854,
775,807

Use for large
integers

Doubles (64bit)

Numbers with
Decimal Places

Decimal values are
approximated

Use for floats where
high precision is not
required

“High Precision
Doubles” (128bit)

Numbers with
Decimal Places

Decimal values are
stored with high
precision (34 decimal
digits)

Use for floats where
high precision is
required

High Precision Floating Point Numbers

Doubles (64bit Floats)

MongoDB Default for ALL
Numbers

Higher Range of Numbers
but lower Decimal Precision

Decimal (128bit Floats)

Has to be Created Explicitly

Lower Range of Numbers but
higher Decimal Precision

Security & User Authentication

Lock Down Your Data

Security Checklist

Authentication &
Authorization

Transport Encryption

Encryption at Rest

Auditing

Server & Network Config
and Setup

Backups & Software
Updates

Authentication & Authorization

Authentication

Identifies valid users of the database

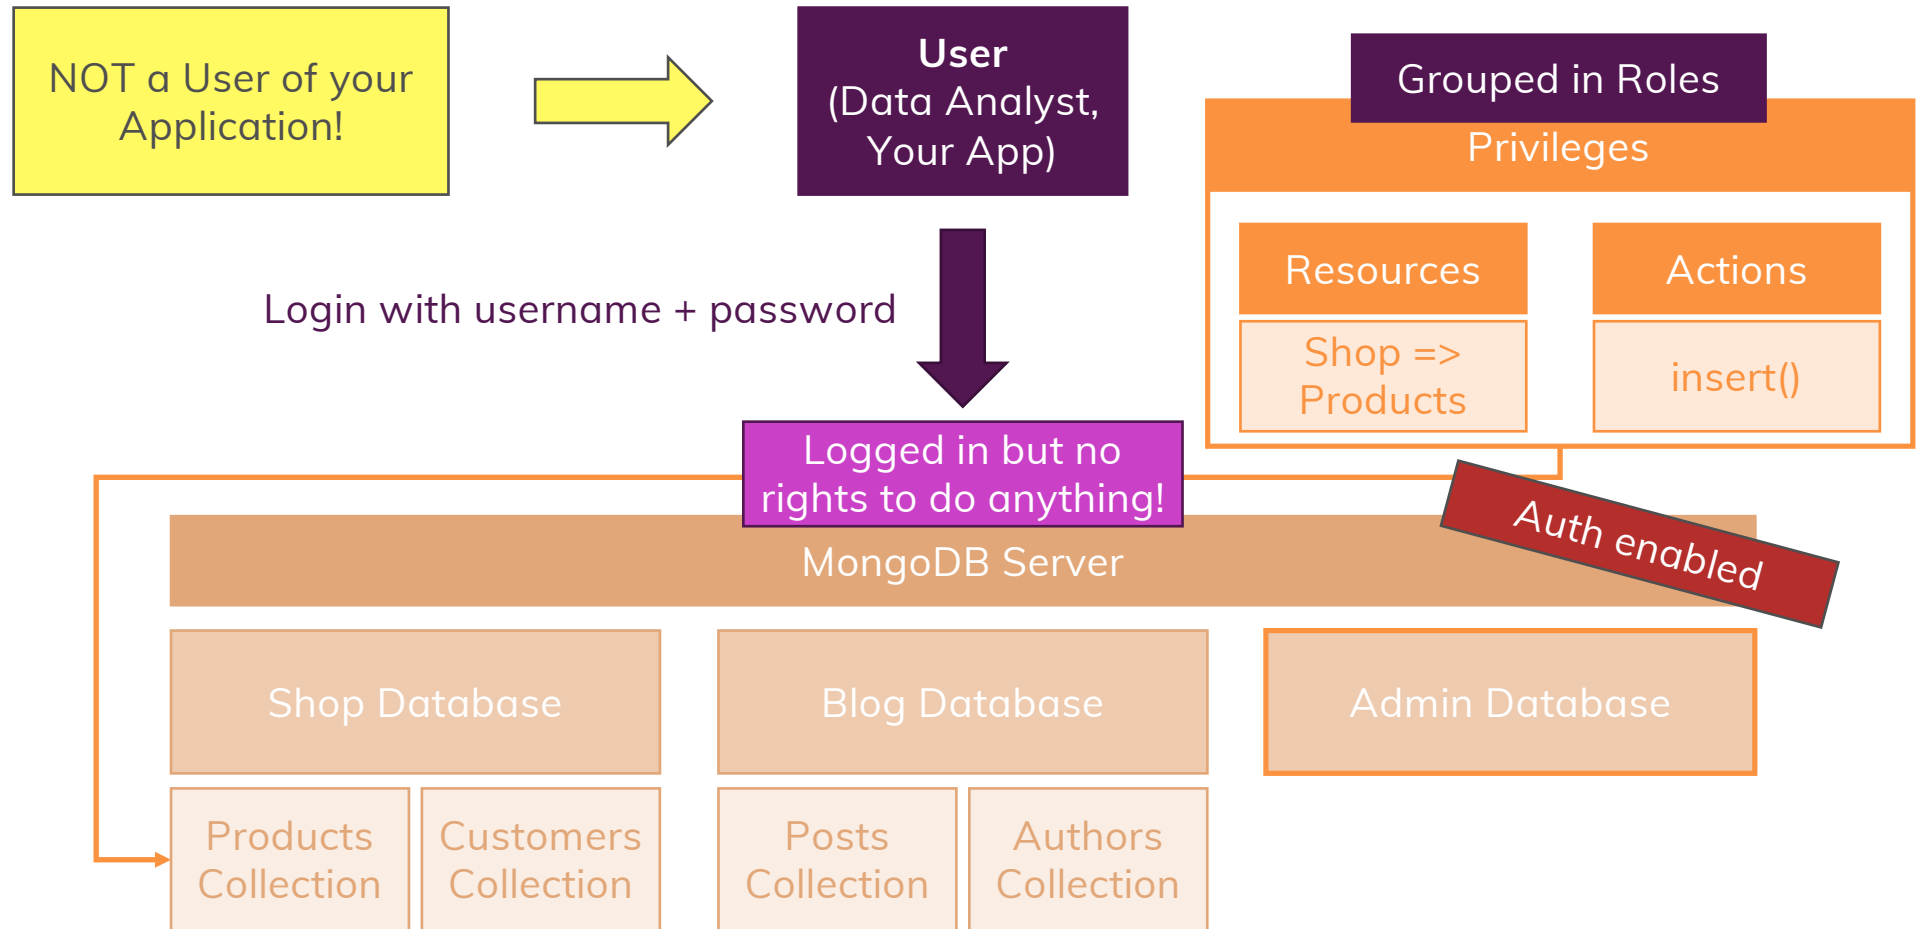
Analogy: You are employed and therefore may access the office

Authorization

Identifies what these users may actually do in the database

Analogy: You are employed as an account and therefore may access the office and process orders

Role Based Access Control



Why Roles?

Different Types of Database Users



Administrator

Needs to be able to manage the database config, create users etc

Does NOT need to be able to insert or fetch data

Developer / Your App

Needs to be able to insert, update, delete or fetch data (CRUD)

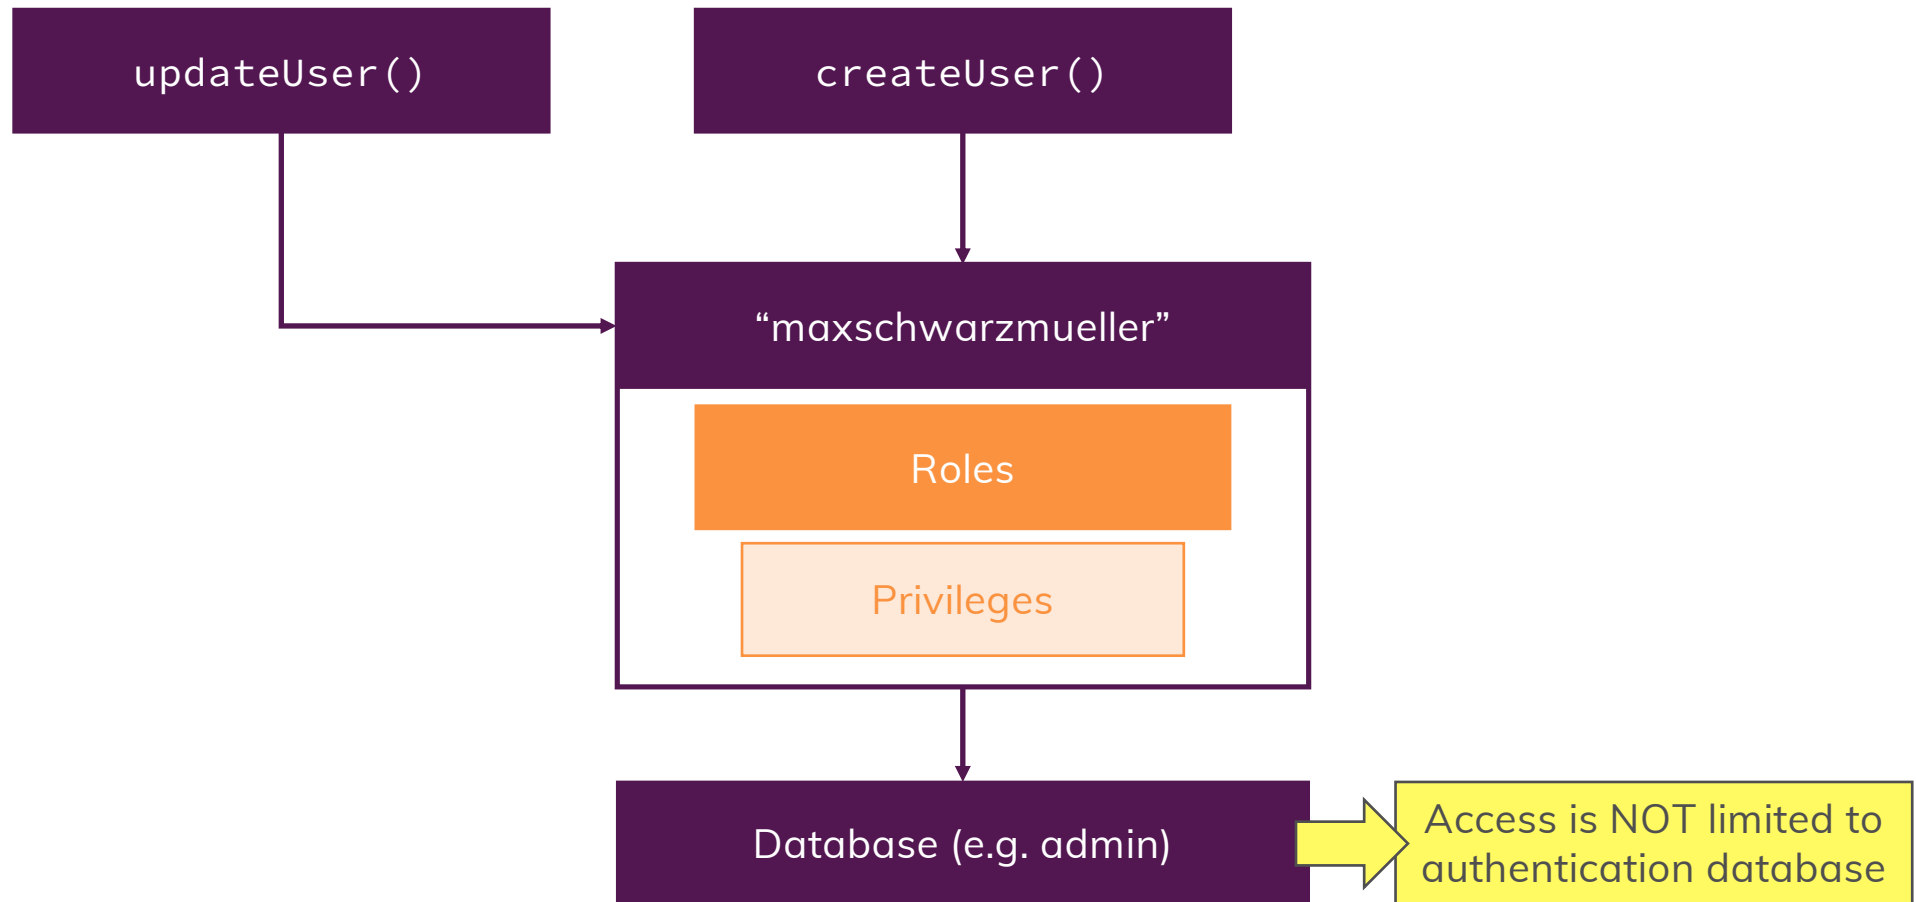
Does NOT need to be able to create users or manage the database config

Data Scientist

Needs to be able to fetch data

Does NOT need to be able to create users, manage the database config or insert, edit or delete data

Creating & Editing Users



Built-in Roles

Database User

read
readWrite

Database Admin

dbAdmin
userAdmin
dbOwner

All Database Roles

readAnyDatabase
readWriteAnyDatabase
userAdminAnyDatabase
dbAdminAnyDatabase

Cluster Admin

clusterManager
clusterMonitor
hostManager
clusterAdmin

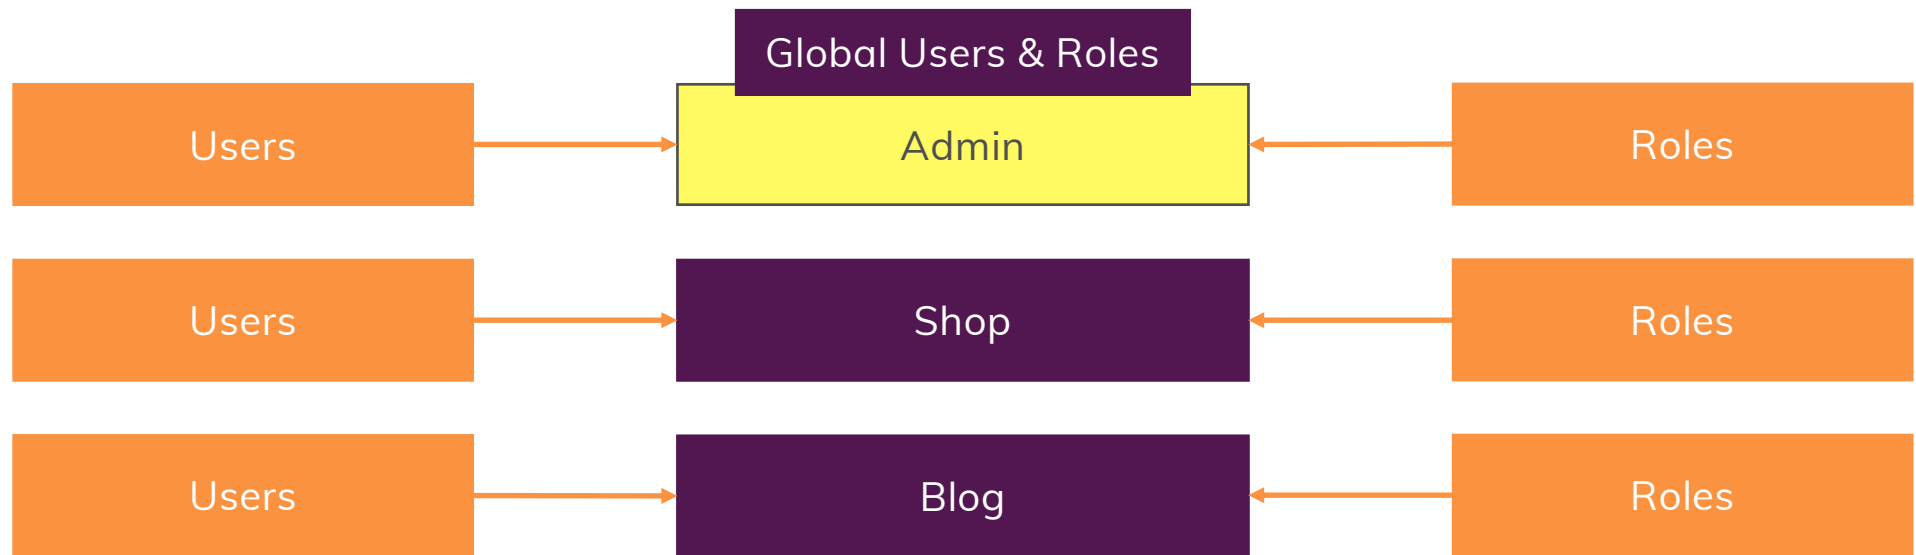
Backup/ Restore

backup
restore

Superuser

dbOwner (admin)
userAdmin (admin)
userAdminAnyDatabase
root

What's Up With The Databases?



User authenticate against their Database

Access is **NOT limited** to that Database though because **Roles define Access Rights**

Roles are attached to Databases and can only be assigned to Users who use this Database as an Authentication Database

Practice!

Database Admin

Work on Database, Create
Collections, Create
Indexes

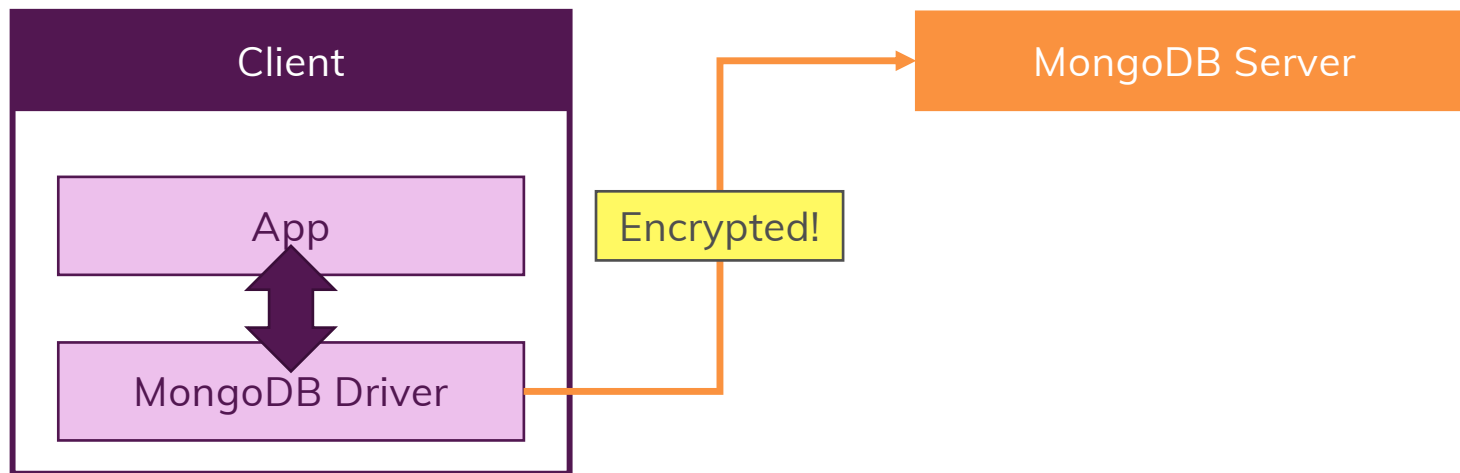
User Admin

Manage Users

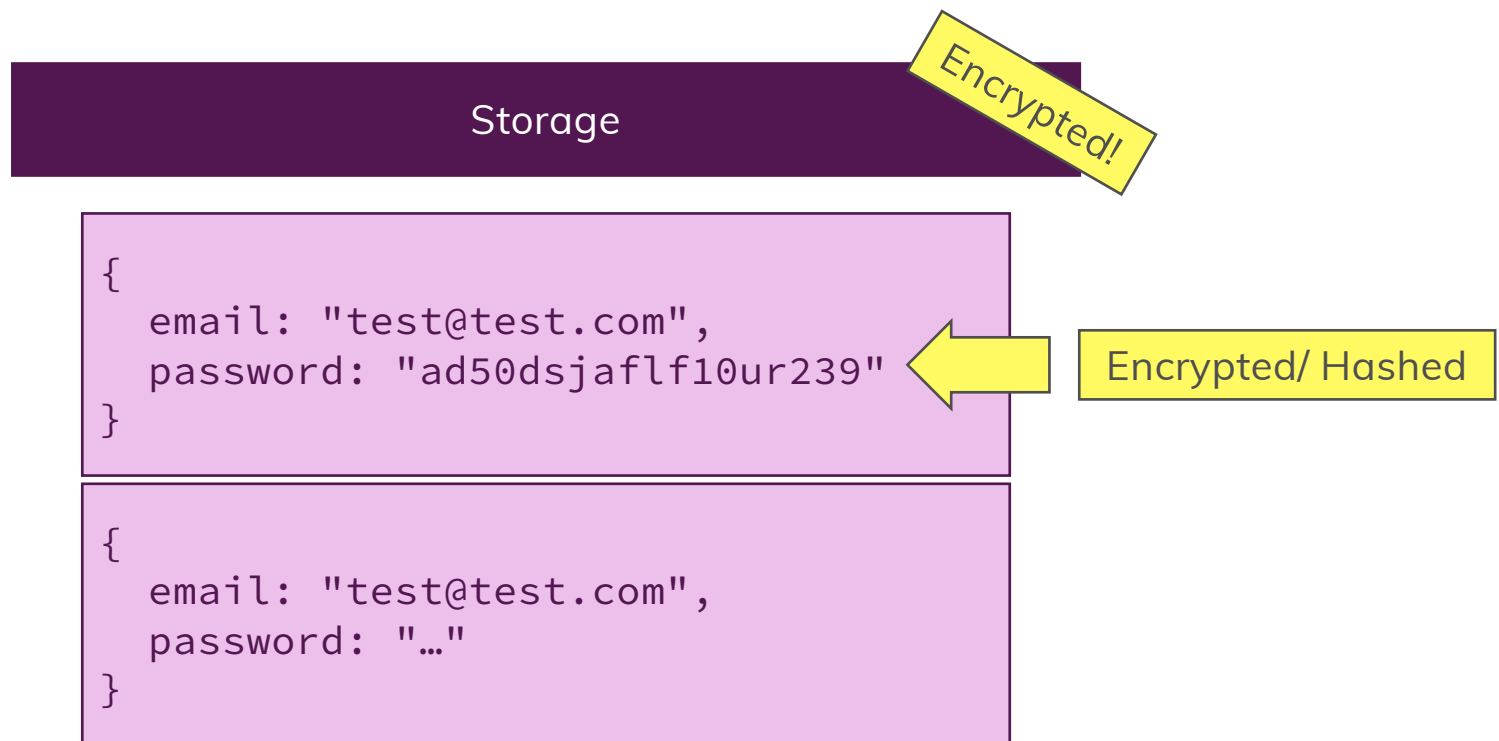
Developer

Read & Write Data in
"Customers" and "Sales"
Databases

Transport Encryption



Encryption at Rest



Module Summary

Users & Roles

- MongoDB uses a Role Based Access Control approach
- You create users on databases and you then log in with your credentials (against those databases)
- Users have no rights by default, you need to add roles to allow certain operations
- Permissions that are granted by roles (“Privileges”) are only granted for the database the user was added to unless you explicitly grant access to other databases
- You can use “AnyDatabase” roles for cross-database access

Encryption

- You can encrypt data during transportation and at rest
- During transportation, you use TLS/ SSL to encrypt data
- For production, you should use SSL certificates issues by a certificate authority (NOT self-signed certificates)
- For encryption at rest, you can encrypt both the files that hold your data (made simple with “MongoDB Enterprise”) and the values inside your documents

Performance, Fault Tolerance & Deployment

Entering the Enterprise World

What's Inside This Module?

What influences Performance?


Capped Collections

Replica Sets


Sharding

MongoDB Server Deployment


What Influences Performance?



Replica Sets



Replica Sets Reads

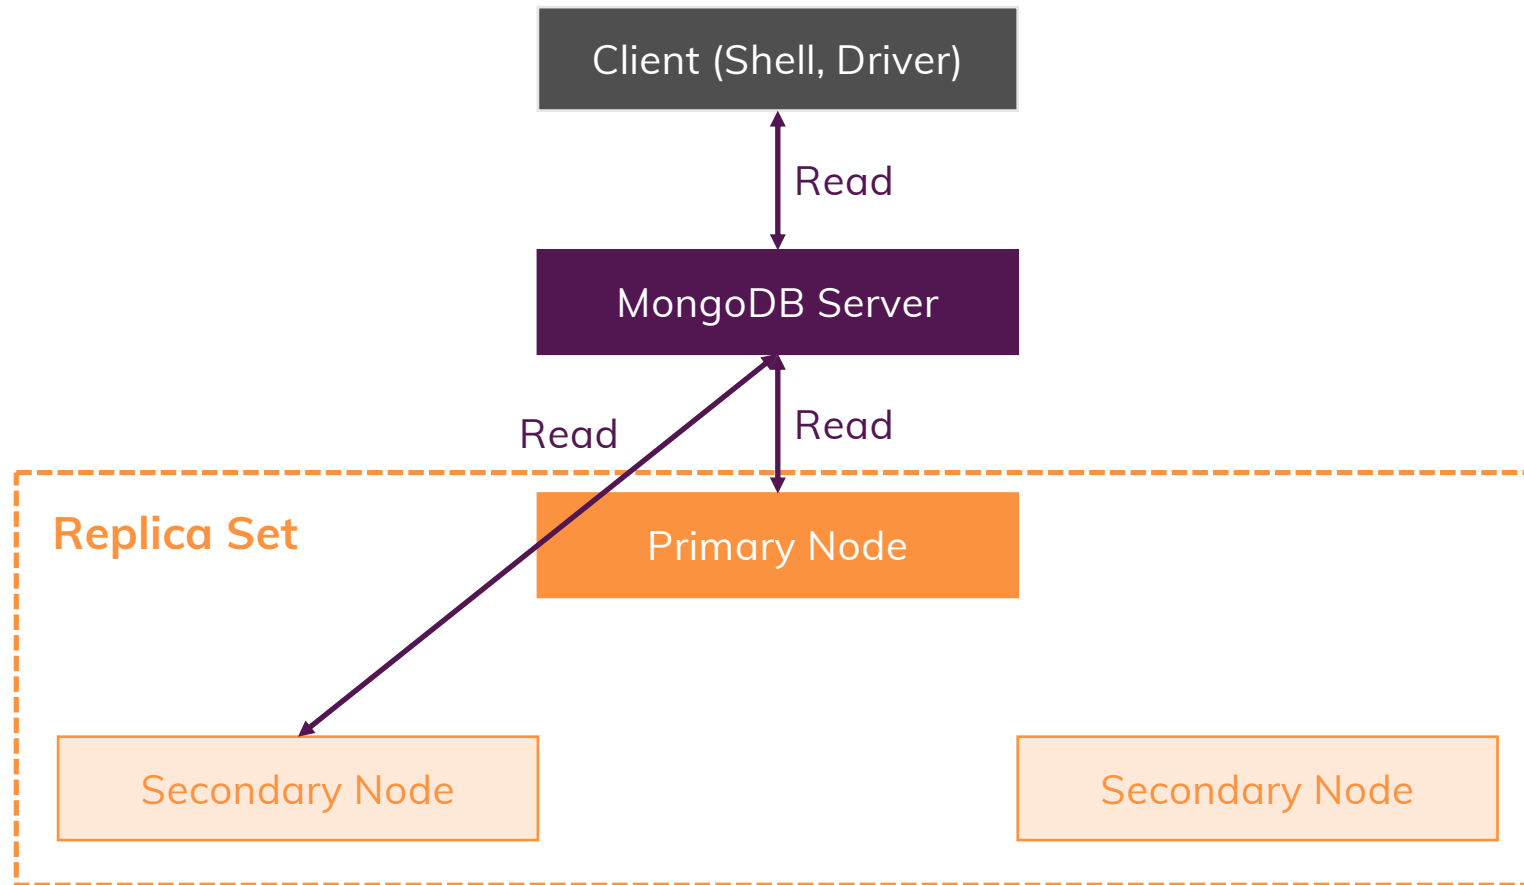


Why Replica Sets?

Backup / Fault Tolerancy

Improve Read Performance

Replica Sets Secondary Reads



Sharding (Horizontal Scaling)

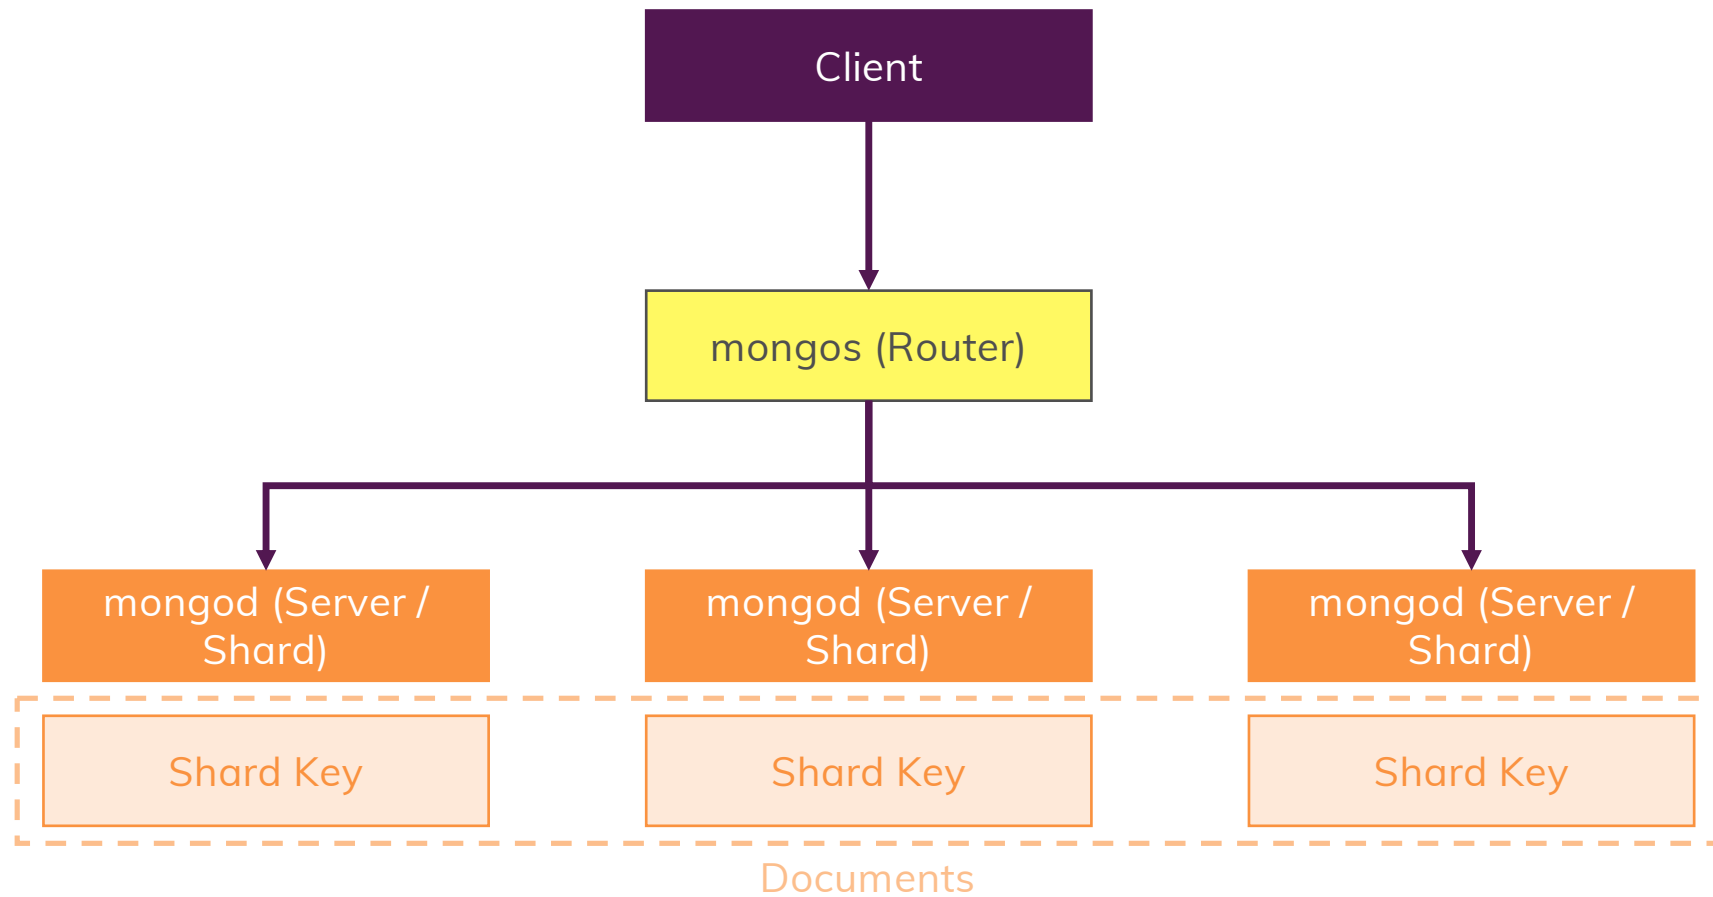
MongoDB Server



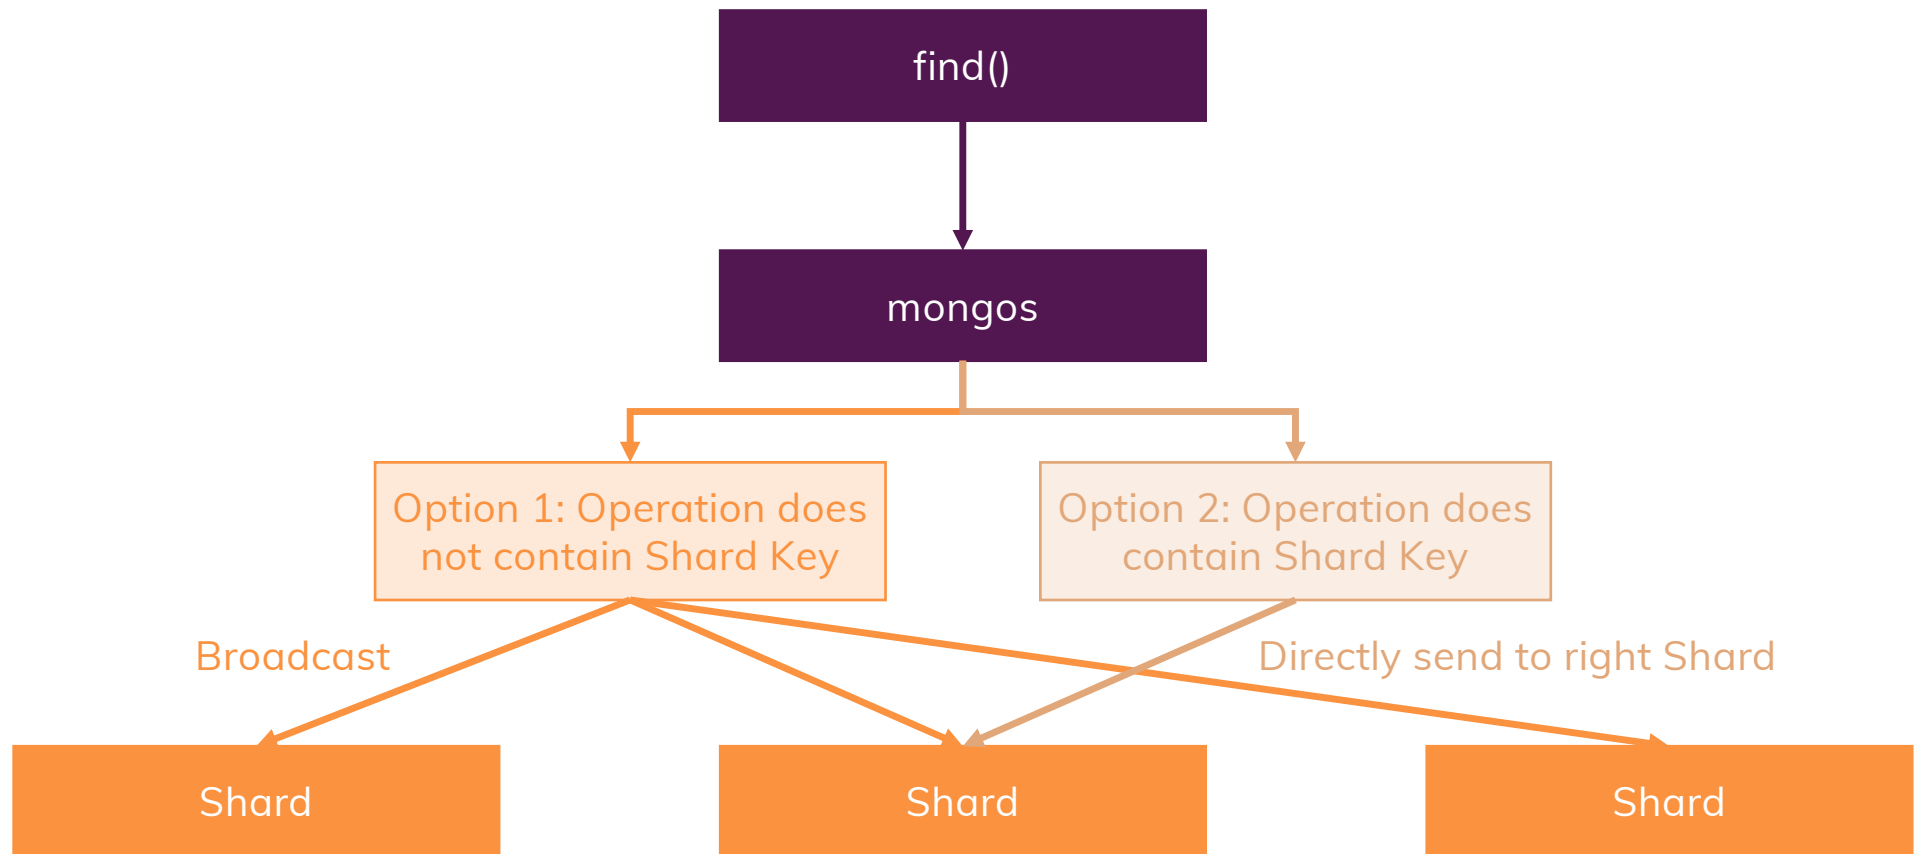
Data is distributed (not replicated!) across Shards

Queries are run across all Shards


How Sharding Works



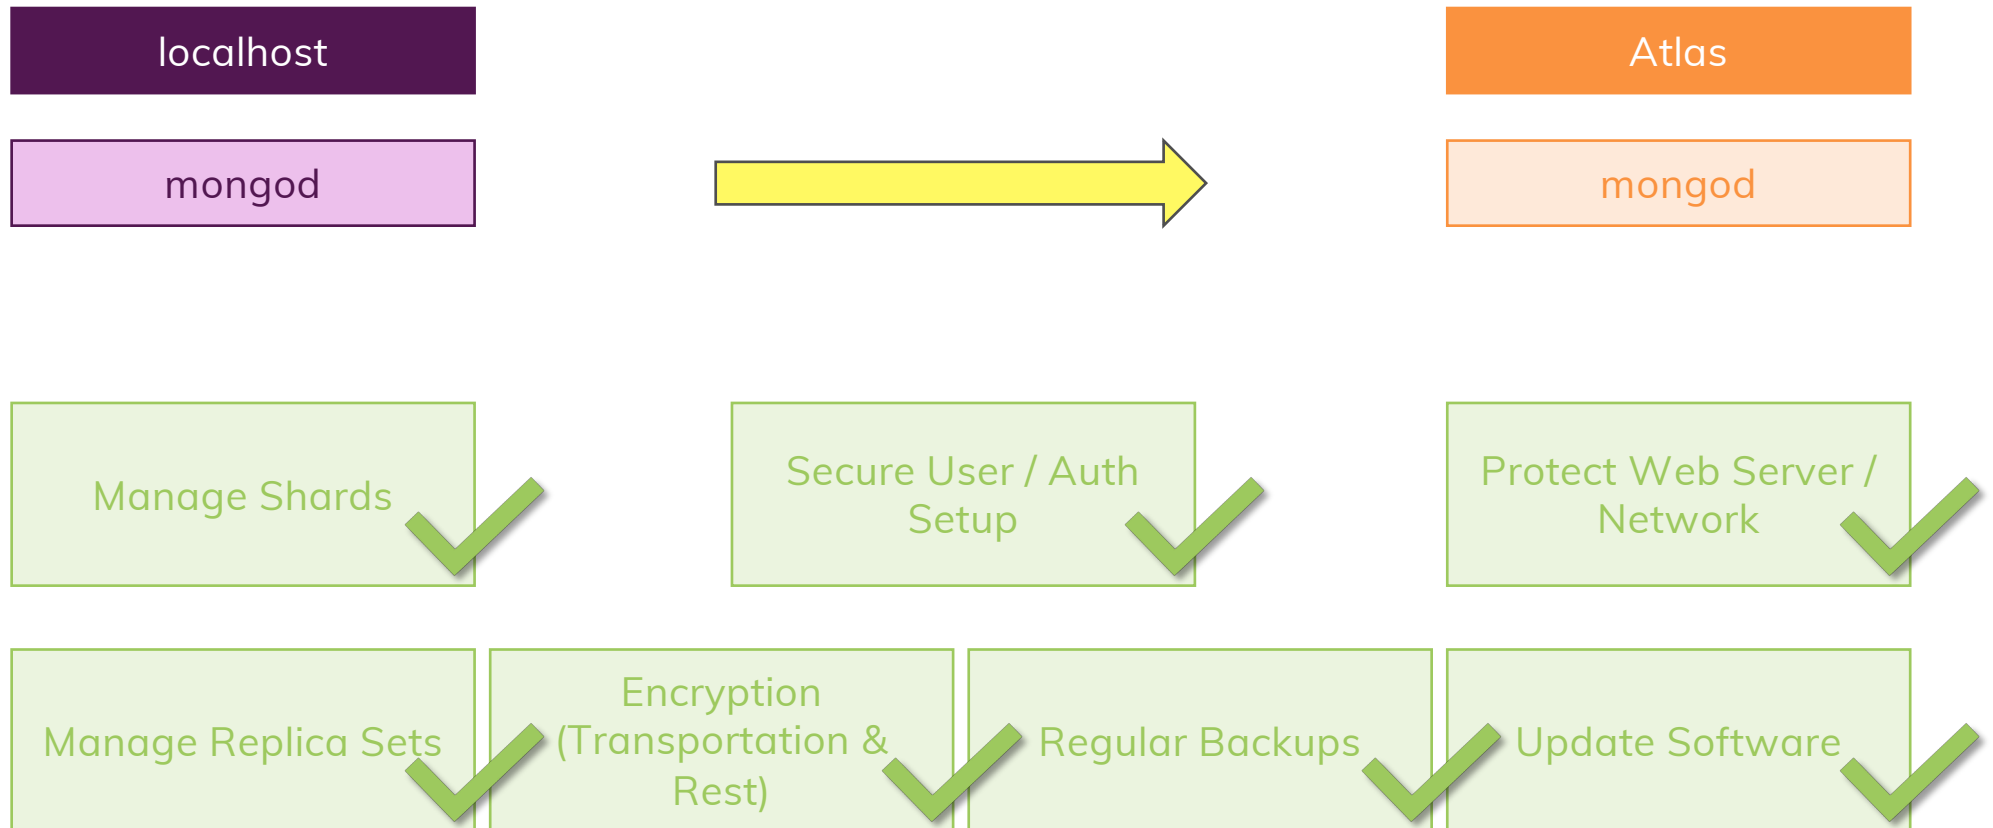
Queries & Sharding



Deploying a MongoDB Server



MongoDB Atlas is a Managed Solution



Module Summary

Performance & Fault Tolerancy

- Consider Capped Collections for cases where you want to clear old data automatically
- Performance is all about having efficient queries/ operations, fitting data formats and a best-practice MongoDB server config
- Replica sets provide fault tolerancy (with automatic recovery) and improved read performance
- Sharding allows you to scale your MongoDB server horizontally

Deployment & MongoDB Atlas

- Deployment is a complex matter since it involves many tasks – some of them are not even directly related to MongoDB
- Unless you are an experienced admin (or you got one), you should consider a managed solution like MongoDB Atlas
- Atlas is a managed service where you can configure a MongoDB environment and pay at a by-usage basis

Transactions

Fail Together

Transactions

User deletes Account

Users Collection

Posts Collection


{ User Document }

related

{ Post Document }

{ Post Document }

Should be deleted together



From Mongo Shell to Drivers

Writing Application Code

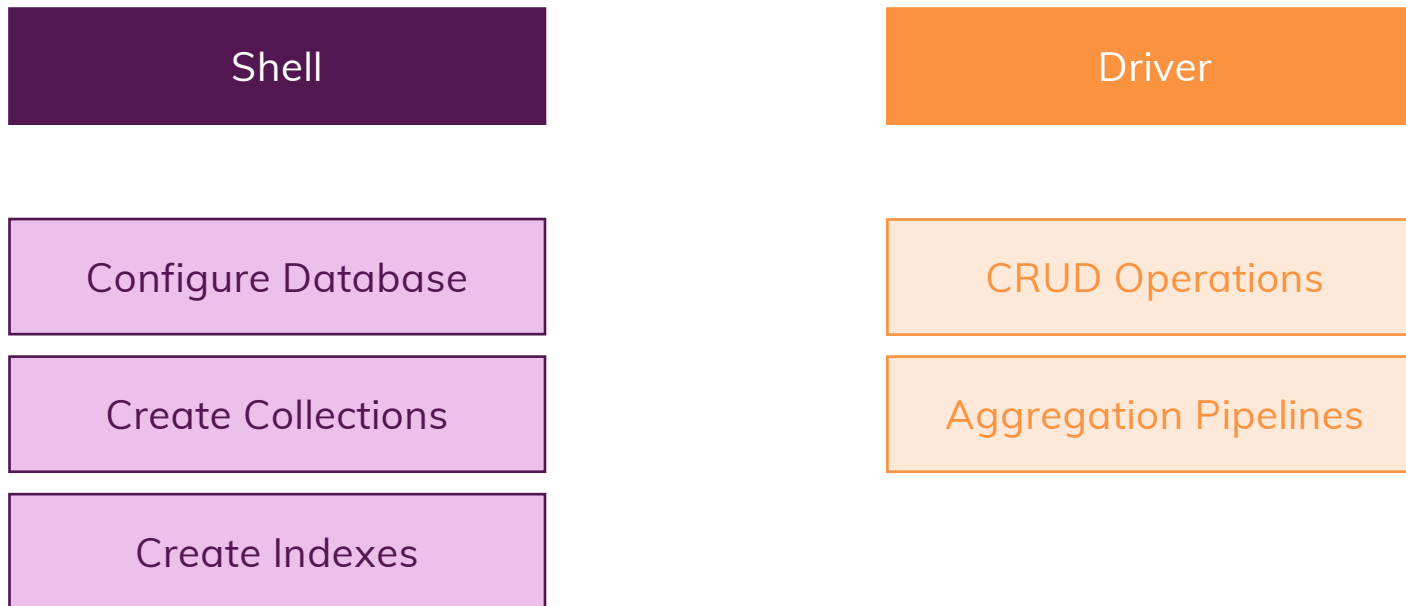
What's Inside This Module?

How translate “Shell Commands” to
“Driver Commands”

Connecting to MongoDB Servers

CRUD Operations

Splitting Work between Drivers & Shell



MongoDB Stitch

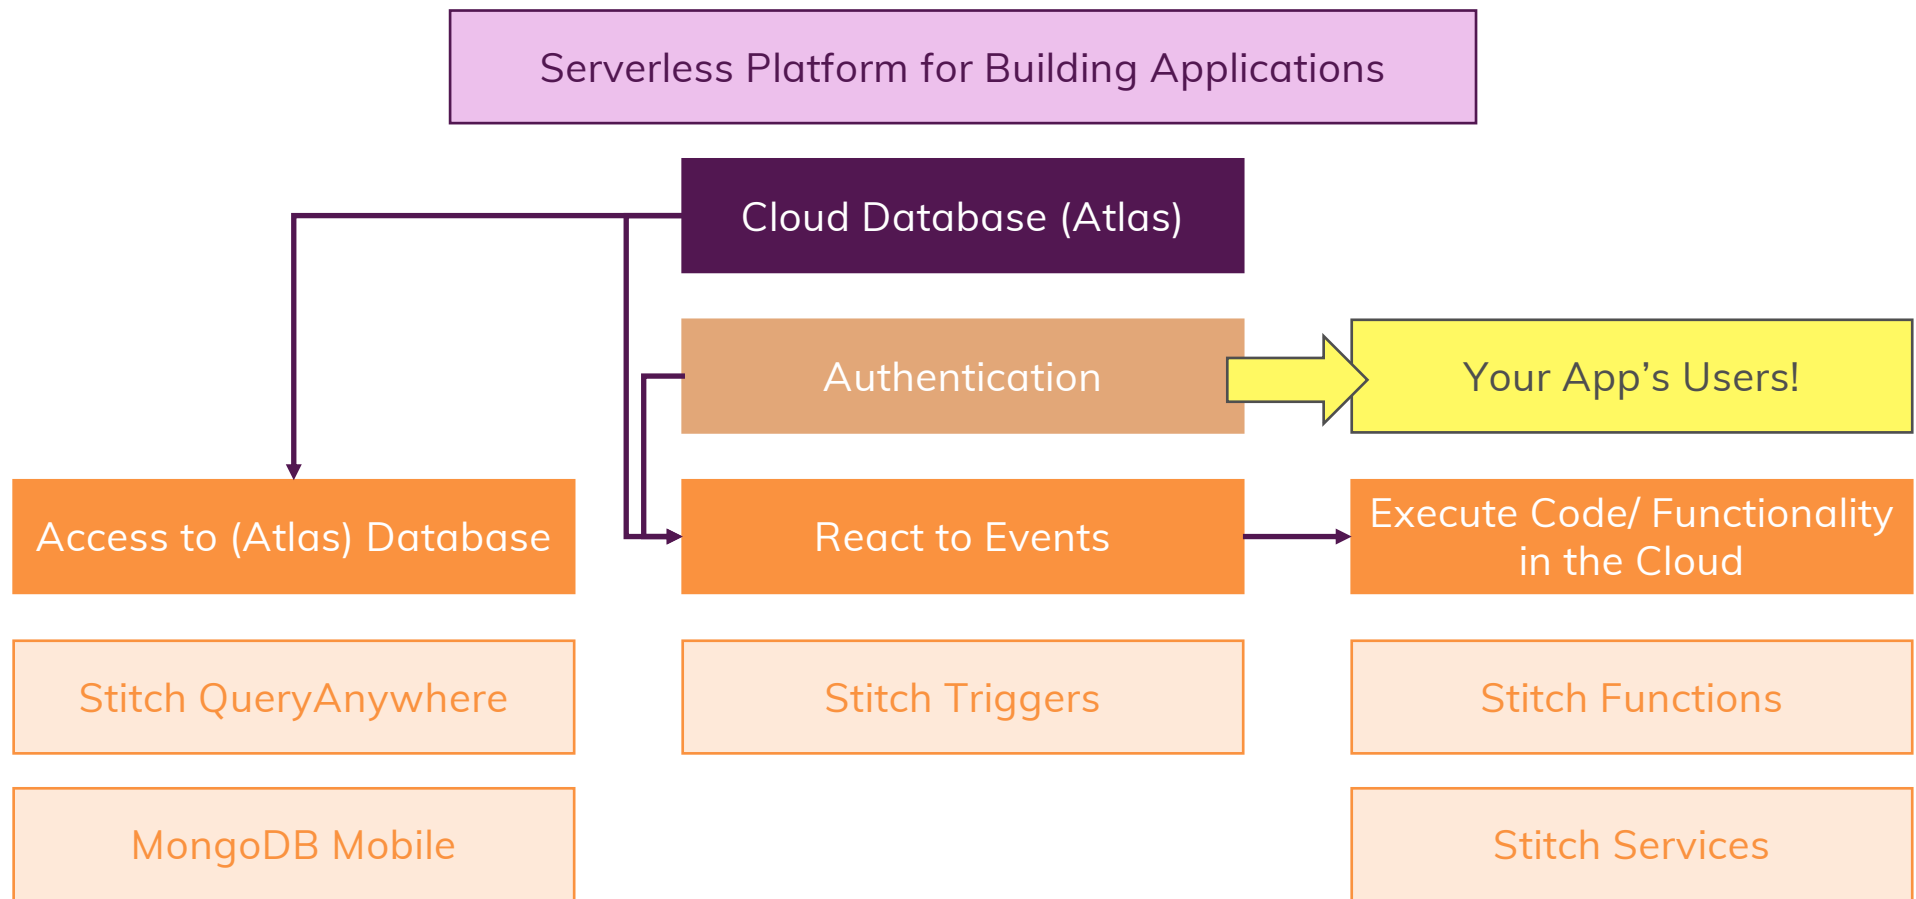
Beyond Data Storage

What's Inside This Module?


What is Stitch?

Using Stitch

What is Stitch?



Serverless?



Stitch Authentication vs MongoDB Authentication

Stitch Authentication

MongoDB stores + manages your Application Users

Signup + Login via Stitch SDK

No Credentials have to be exposed in Clients

Highly Granular Permissions

MongoDB Authentication

Your create + manage Database Users

Login during Connection

MongoDB Credentials have to be exposed => Not usable in Clients

Role-based Permissions

Roundup & Next Steps

What Next?

Play Around!

Practice, Practice, Practice

Use the Shell as a
Playground

Build Dummy/ Demo Apps
that use MongoDB

Build Dummy/ Demo Apps
that use MongoDB +
Stitch

Dive into the Official Docs

Dive into Stackoverflow +
Blog Posts (Google!) to
learn Best Practices

Use YouTube + Other
Courses to Learn more
about Specific Topics

Resources